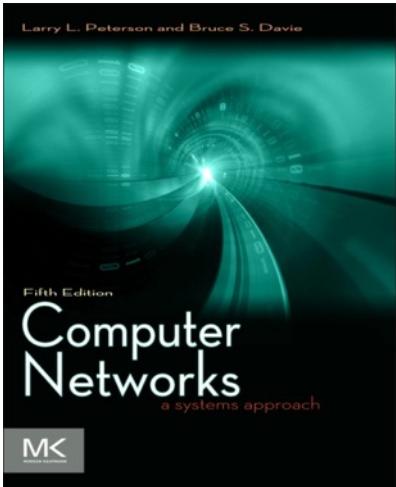
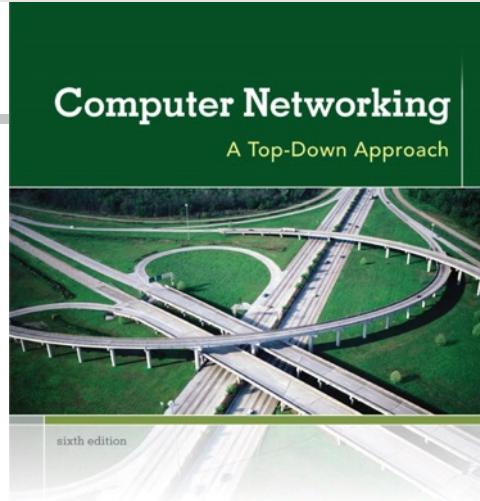


Slides are an edited mashup of two books



Computer Networks:
A Systems Approach, 5e
Larry L. Peterson and Bruce S.
Davie



*Computer Networking: A
Top Down Approach*
6th edition
Jim Kurose, Keith Ross
Addison-Wesley
March 2012

Why is Internet Architecture relevant to you

- The Internet is our economies most critical infrastructure.
- The Internet is our economies most enabling metaphor.
- Objectives for student learning:
 - What are the technical features that allowed the Internet to become so pervasive and diverse during the course of your lifetime.
 - Understand its basic structure, and distill from it lessons for
 - (a) the use and governance of Internet technology in modern organizations, and
 - (b) lessons for creation of other new technologies and products.

Internet history

1972-1980: Internetworking, new and proprietary nets

- 1970: ALOHAnet satellite network in Hawaii
- 1974: Cerf and Kahn - architecture for interconnecting networks
- 1976: Ethernet at Xerox PARC
- late70' s: proprietary architectures: DECnet, SNA, XNA
- late 70' s: switching fixed length packets (ATM precursor)
- 1979: ARPAnet has 200 nodes

Cerf and Kahn's internetworking principles:

- minimalism, autonomy - no internal changes required to interconnect networks
- best effort service model
- stateless routers
- decentralized control

define today's Internet architecture

Internet history

1980-1990: new protocols, a proliferation of networks

- 1983: deployment of TCP/IP
- 1982: smtp e-mail protocol defined
- 1983: DNS defined for name-to-IP-address translation
- 1985: ftp protocol defined
- 1988: TCP congestion control
- new national networks: Csnet, BITnet, NSFnet, Minitel
- 100,000 hosts connected to confederation of networks

Internet history

1990, 2000's: commercialization, the Web, new apps

- early 1990's: ARPAnet decommissioned
- 1991: NSF lifts restrictions on commercial use of NSFnet (decommissioned, 1995)
- early 1990s: Web
 - hypertext [Bush 1945, Nelson 1960's]
 - HTML, HTTP: Berners-Lee
 - 1994: Mosaic, later Netscape
 - late 1990's: commercialization of the Web
- late 1990's – 2000's:
 - more killer apps: instant messaging, P2P file sharing
 - network security to forefront
 - est. 50 million host, 100 million+ users
 - backbone links running at Gbps

Internet history

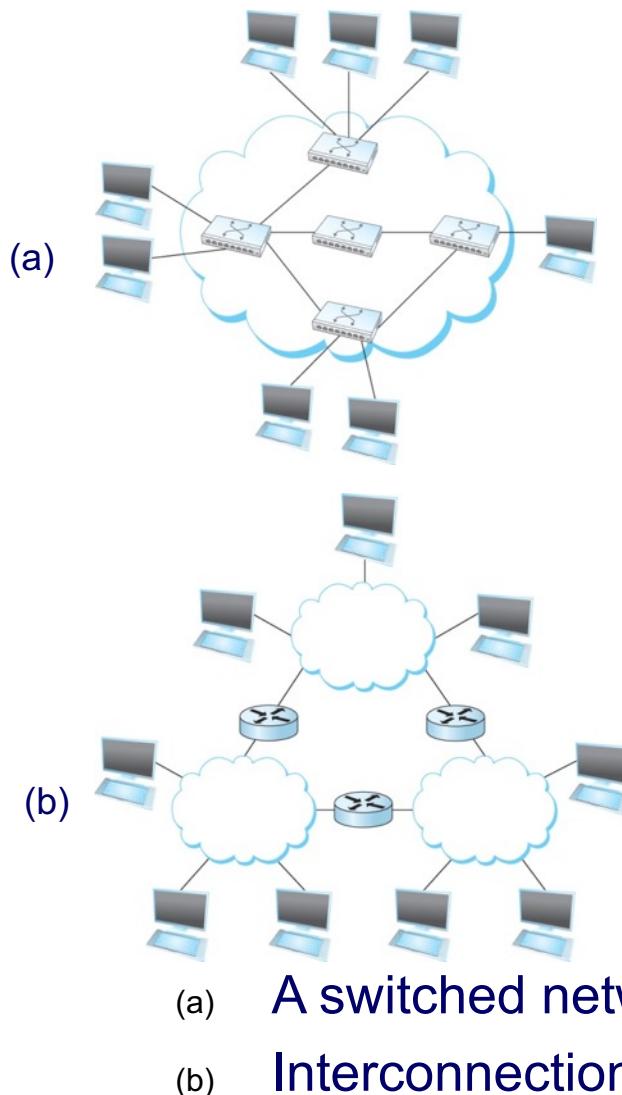
2005-present

- ~750 million hosts
 - Smartphones and tablets
- Aggressive deployment of broadband access
- Increasing ubiquity of high-speed wireless access
- Emergence of online social networks:
 - Facebook, Instagram, ...
- Service providers (Google, Amazon) create their own networks
 - Bypass Internet, providing “instantaneous” access to search, email, etc.
- E-commerce, universities, enterprises running their services in “cloud” (eg, Amazon EC2)

What makes the Internet appear as single service

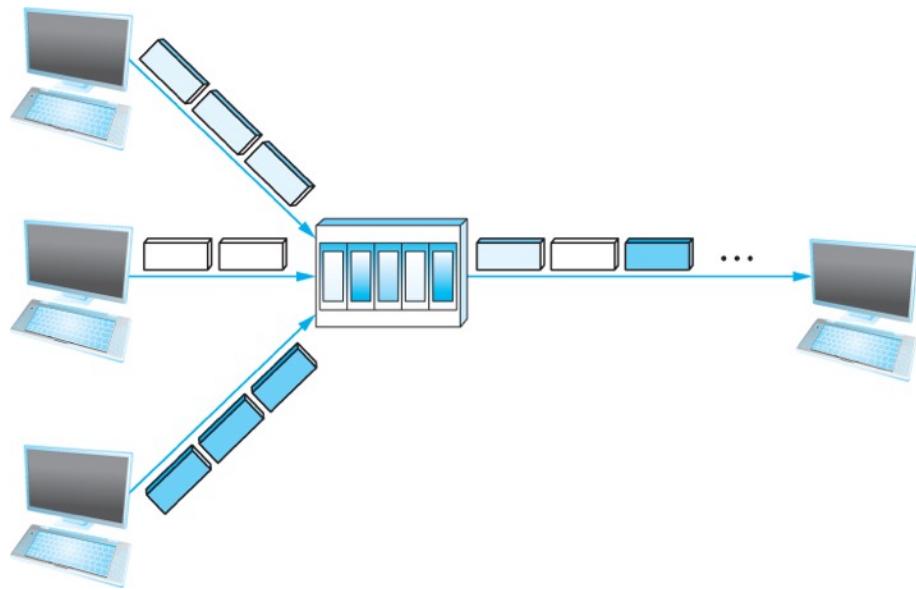
- Networks share common architecture and protocols that enable communication within and among them.
 - Architecture: how components of the networks interrelate
 - Protocols: standards governing the interchange of data
- The architecture and protocols were and to some extent are shaped by fundamental (and interrelated) design principles adopted by early builders of the Internet.
 - Hourglass
 - End-to-End
 - Distributed design and decentralized control
 - Heterogeneity, Scalability

Connectivity Terminologies



- Link, Nodes
- Point-to-point, Multiple access
- Switched Network
 - Circuit Switched
 - Packet Switched: Store-and-forward
- Cloud
- Hosts
- Switches
- internetwork
- Router/gateway
- Host-to-host connectivity
- Address
- Routing
- Unicast/broadcast/multicast

Cost-Effective Resource Sharing

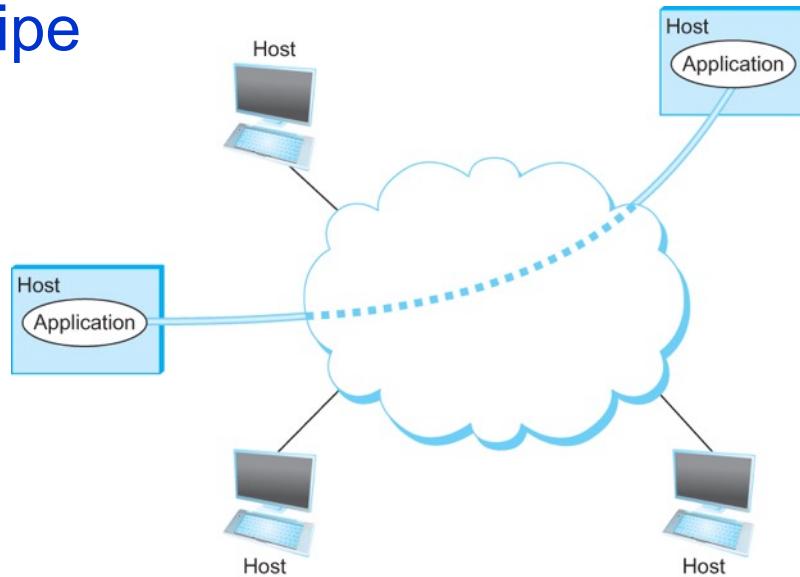


A switch multiplexing packets from multiple sources onto one shared link

- FDM: Frequency Division Multiplexing
- Statistical Multiplexing
 - Data is transmitted based on demand of each flow.
 - What is a flow?
 - Packets vs. Messages
 - FIFO, Round-Robin, Priorities (Quality-of-Service (QoS))
 - Congested?
- LAN, MAN, WAN
- SAN (System Area Networks)

Support for Common Services

- Logical Channels
 - Application-to-Application communication path or a pipe



Process communicating over an abstract channel

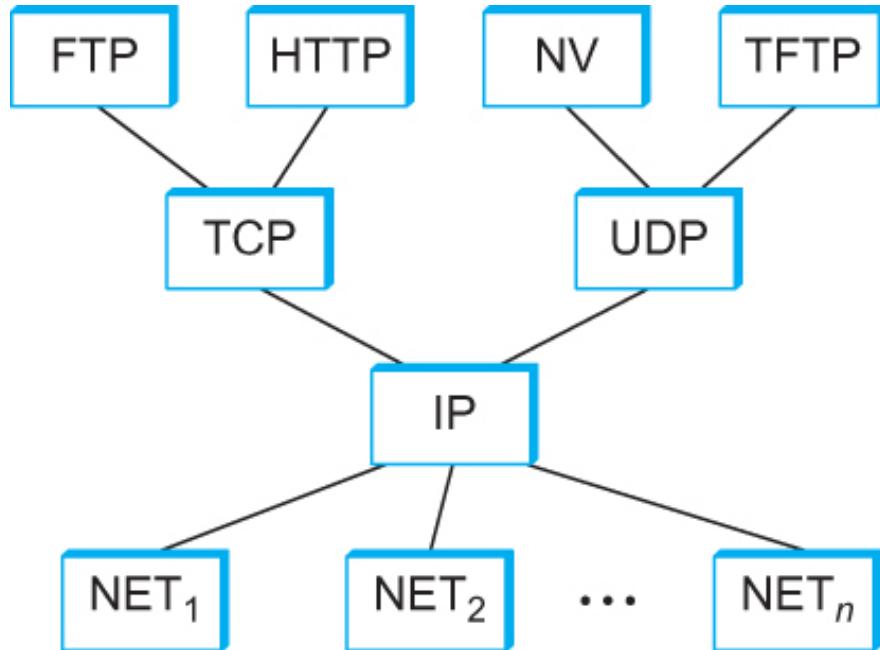
Reliability challenge

- Network software should hide (inevitable) errors from applications
- Bits are lost
 - Bit errors (1 to a 0, and vice versa)
 - Burst errors – several consecutive errors
- Packets are dropped (largely Congestion)
 - Links and Node failures
 - Messages are delayed
- Messages are delivered out-of-order
- Third parties eavesdrop

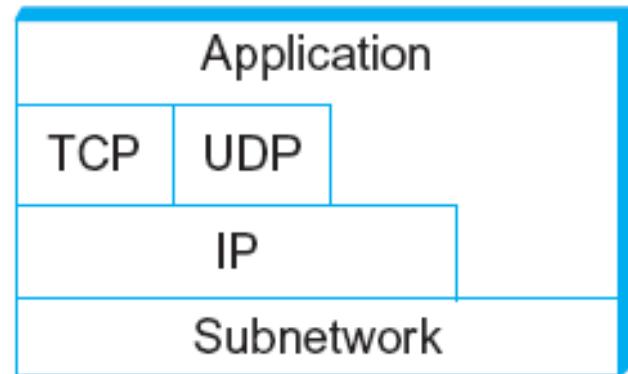
Protocols

- Protocol defines the interfaces between the layers in the same system and with the layers of peer system
- Building blocks of a network architecture
- Each protocol object has two different interfaces
 - service interface: operations on this protocol
 - peer-to-peer interface: messages exchanged with peer
- Term “protocol” is overloaded
 - specification of peer-to-peer interface
 - module that implements this interface

Internet Architecture



Internet Protocol Graph



Alternative view of the Internet architecture. The “Network” layer shown here is sometimes referred to as the “sub-network” or “link” layer.

Description of (Lower) Layers

- Physical Layer
 - Handles the transmission of raw bits over a communication link
- Data Link Layer
 - Collects a stream of bits into a larger aggregate called a *frame*
 - Network adaptor along with device driver in OS implement the protocol in this layer
 - Frames are actually delivered to hosts
- Network Layer
 - Handles routing among nodes within a packet-switched network
 - Unit of data exchanged between nodes in this layer is called a *packet*

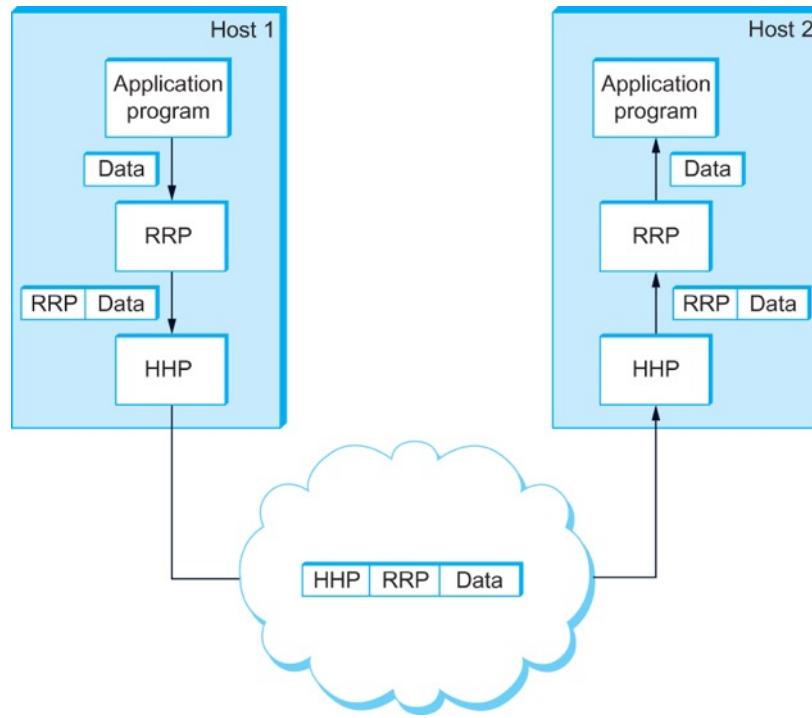
The lower three layers are implemented on all network nodes

Description of (Higher) Layers

- Transport Layer
 - Implements a process-to-process channel
 - Unit of data exchanges in this layer is called a *message*
- Session Layer
 - Provides a name space that is used to tie together the potentially different transport streams that are part of a single application
- Presentation Layer
 - Concerned about the format of data exchanged between peers
- Application Layer
 - Standardize common type of exchanges

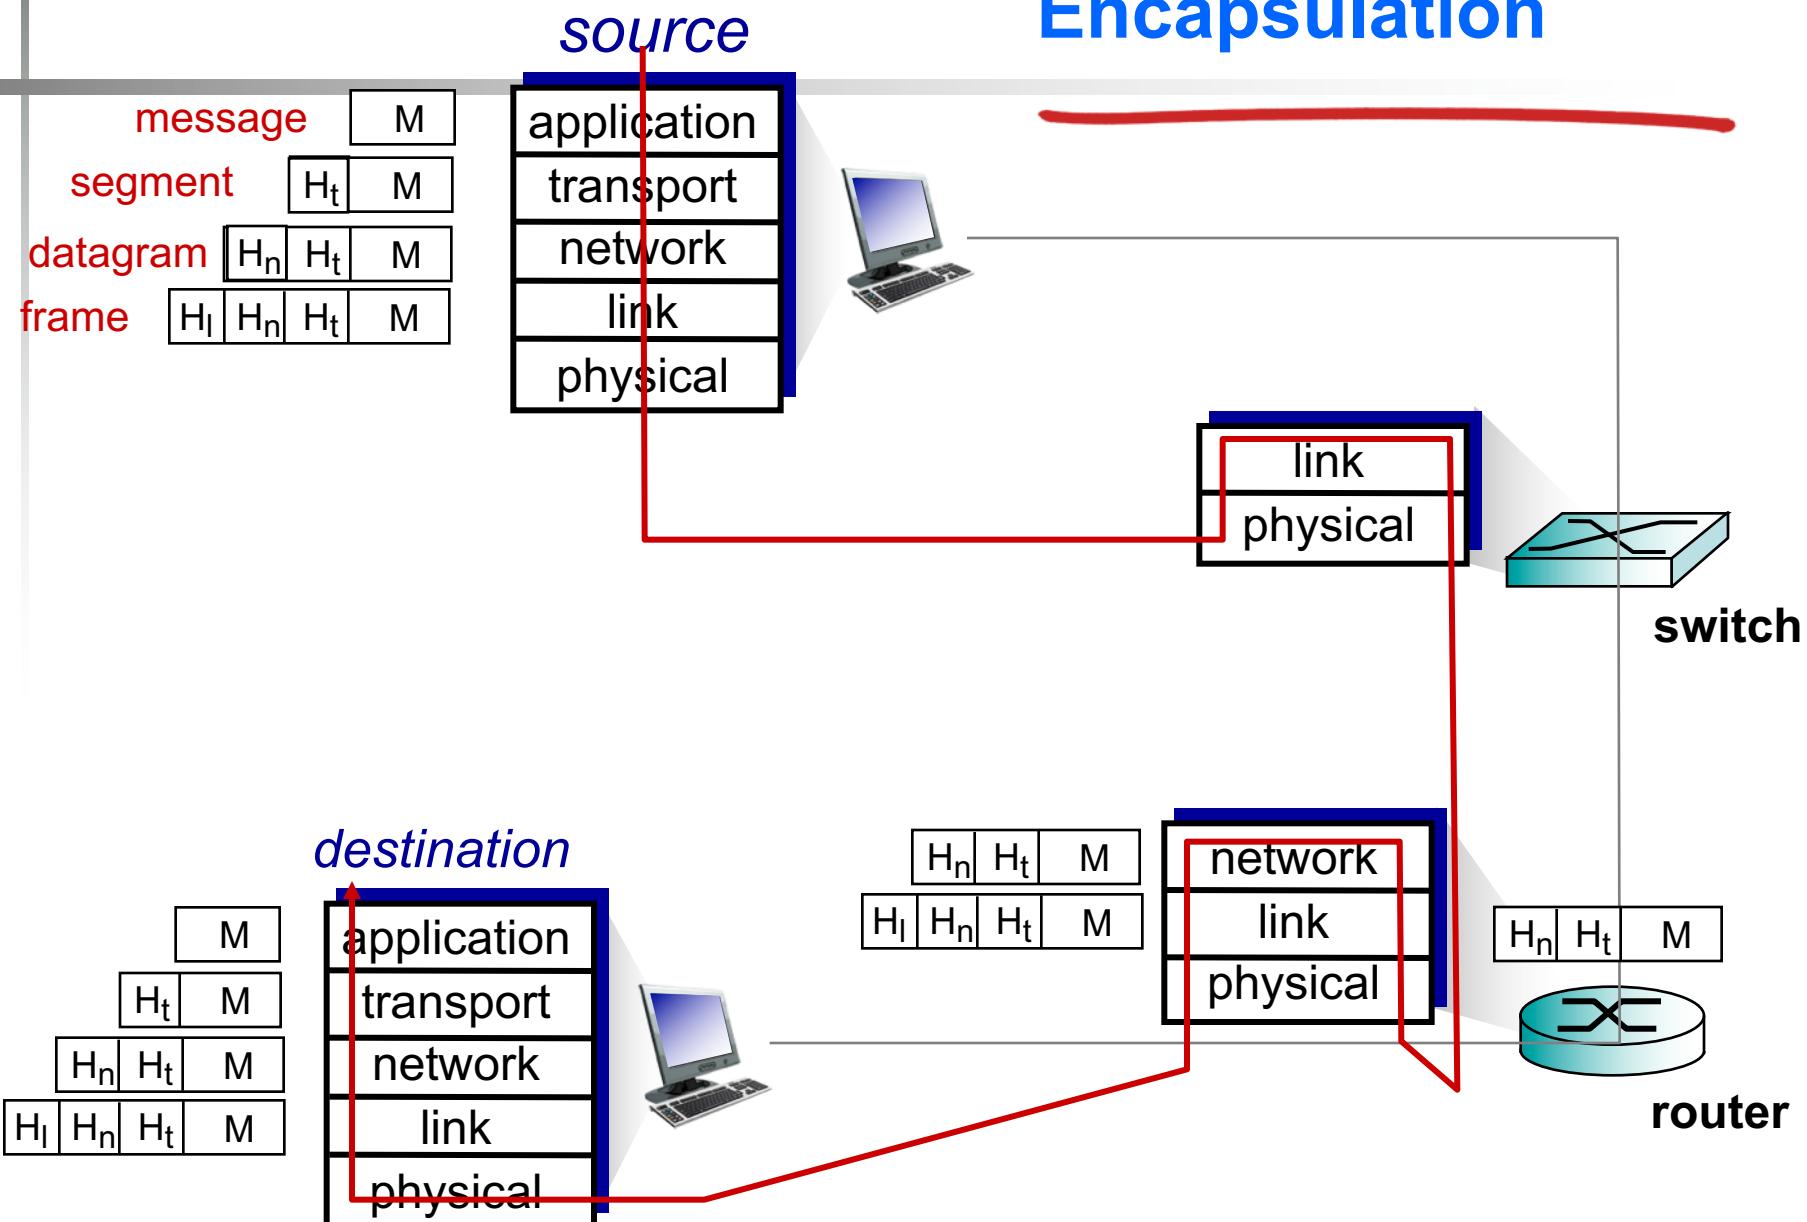
The transport layer and the higher layers typically run only on end-hosts and not on the intermediate switches and routers

Encapsulation



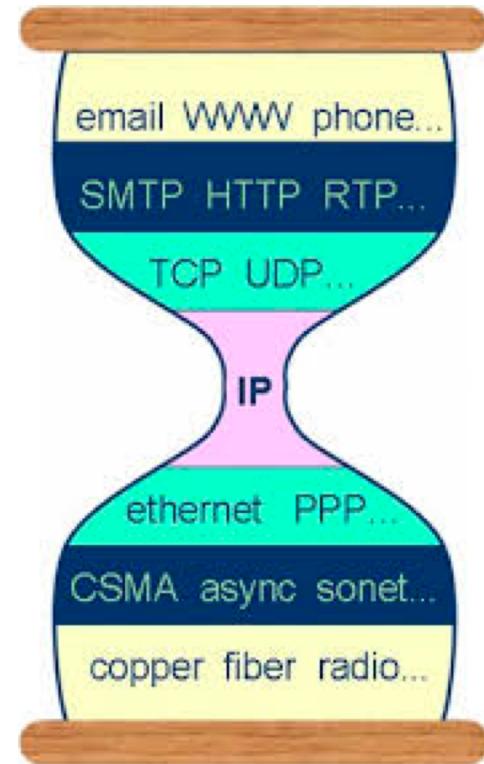
High-level messages are encapsulated inside of low-level messages

Encapsulation



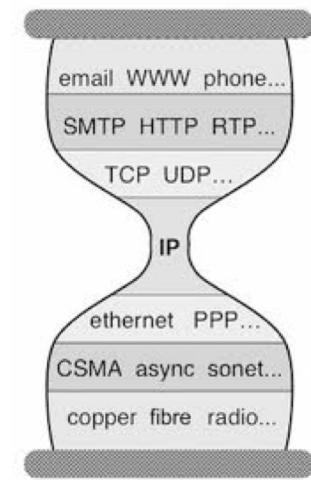
Internet Architecture

- Does not imply strict layering. The application is free to bypass the defined transport layers and to directly use IP or other underlying networks
- An hour-glass shape – wide at the top, narrow in the middle and wide at the bottom. IP serves as the focal point for the architecture
- In order for a new protocol to be officially included in the architecture, there needs to be both a protocol specification and at least one (and preferably two) representative implementations of the specification
- IETF Governance
 - “rough consensus and running code”



Benefits of Hourglass architecture

- Internet designed to operate over different underlying communications technologies, including those yet to be introduced, and to support multiple and evolving applications and services.
- Does not impede or restrict particular applications (although users, ISPs may make optimizations)
- Enables developers to write applications without knowing/adapting to details of underlying networks
- Enables users to adopt applications without involvement/approval from network operators
- Critical separation between network technology and higher-level services through which users actually interact with the Internet visualized as hourglass
- IP as minimal viable agreement/min common denominator maximizes flexibility



Why a narrow waist is important metaphor for new systems, products ...

- Tim O'Reilly:
 - do as little as possible....the less you include the easier it will be to agree and you dont tie yourself down...because we dont know what will come [sic: in this case less is more]
 - Build a system and let it evolve
 - Create architecture for participation—iTUNES, App Store...[sic: it started with Internet, Includes maps mashups and APIs!!]
 - TBL didnt have to ask anyones permission to put up WWW on the net...they would have said no...'http is poorly designed protocol..will never scale'
 - Tolerate as much failure and participation as needed to introduce new systems/innovations rapidly/iteratively and innovate
- Naughton:
 - Allow innovation to be tried for free

End-to-end architecture

- Edge-based innovation derives from early design decision that the Internet should have an end-to- end architecture:
 - The network provides communications fabric connecting the many computers at its ends
 - Network offers very basic level of service, data transport
 - Beyond transporting data—locate special features needed to support specific applications in or close to applications/devices at network edge.
 - Only put feature lower down if performance improvement justifies it
- E2E design facilitates
 - designing for: failure, change, dynamics, decentralized control, rolling asynchronous adoption, of components
 - scalability and therefore longevity of architecture
- QUESTION – Reliable transport of data – Packet level hop by hop, Packet level end to end (Process), Message/File level end to end (Application).

Scalability

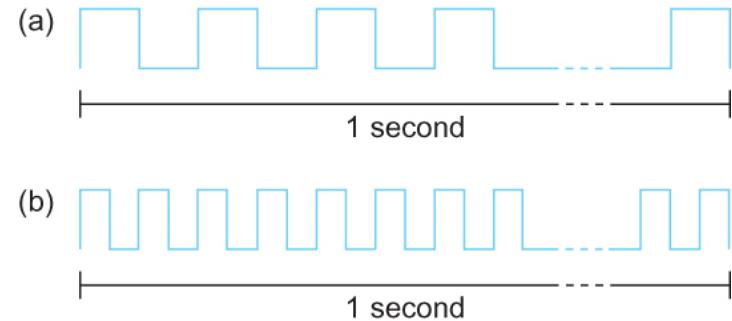
- Internet's design enabled it to support a growing amount of communications:
- Growth in number of users and attached devices
- Growth in volume of communications per device and total
- Scale implies... heterogeneity...designing for Heterogeneity is a good step in future proofing

Performance

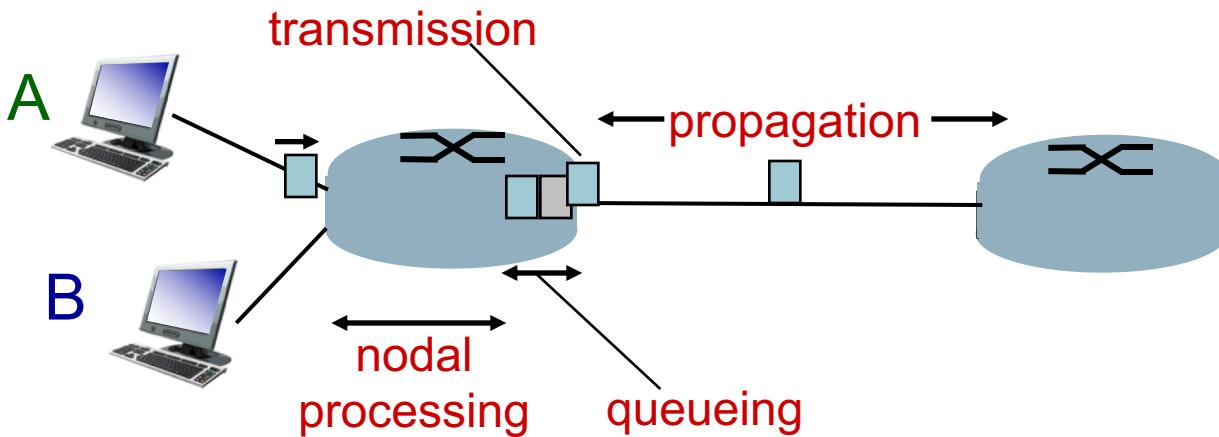
- Latency = Propagation + transmit + queue
 - Propagation = distance/speed of light
 - Transmit = size/bandwidth
-
- One bit transmission => propagation is important
 - Large bytes transmission => bandwidth is important

Bandwidth

- Width of the frequency band
 - Number of bits per second that can be transmitted over a communication link
- 1 Mbps: 1×10^6 bits/second = 1×2^{20} bits/sec
- 1×10^{-6} seconds to transmit each bit or imagine that a timeline, now each bit occupies 1 micro second space.
- On a 2 Mbps link the width is 0.5 micro second.
- Smaller the width more will be transmission per unit time.



Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

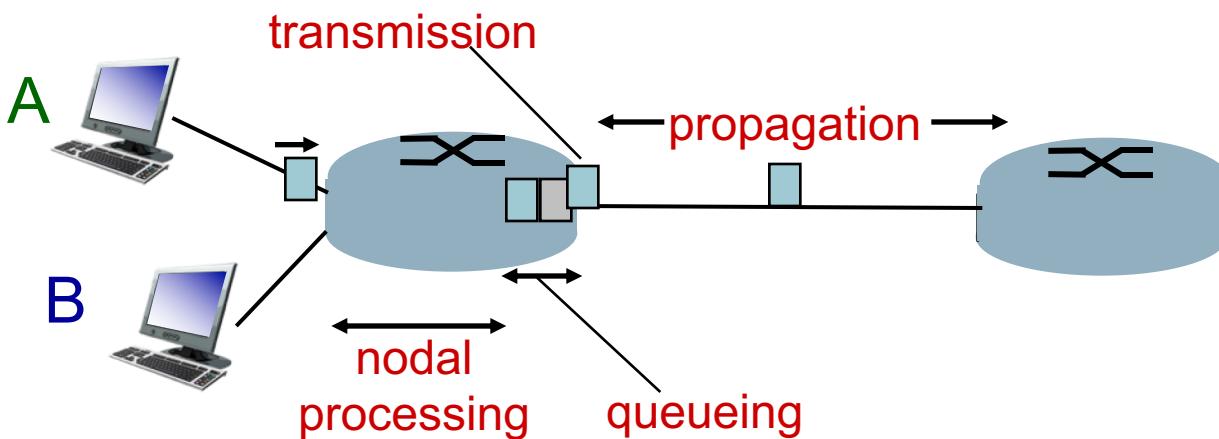
(1) d_{proc} : nodal processing

- check bit errors
- determine output link
- typically < msec

(2) d_{queue} : queueing delay

- time waiting at output link for transmission
- depends on congestion level of router

Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

(3) d_{trans} : transmission delay:

- L : packet length (bits)
- R : link bandwidth (bps)
- $d_{\text{trans}} = L/R$

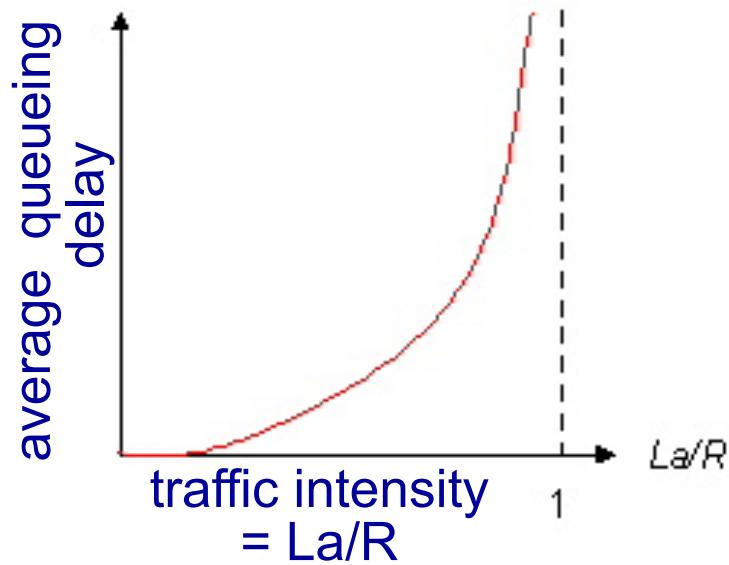
d_{trans} and d_{prop}
very different

(4) d_{prop} : propagation delay:

- d : length of physical link
- s : propagation speed in medium ($\sim 2 \times 10^8 \text{ m/sec}$)
- $d_{\text{prop}} = d/s$

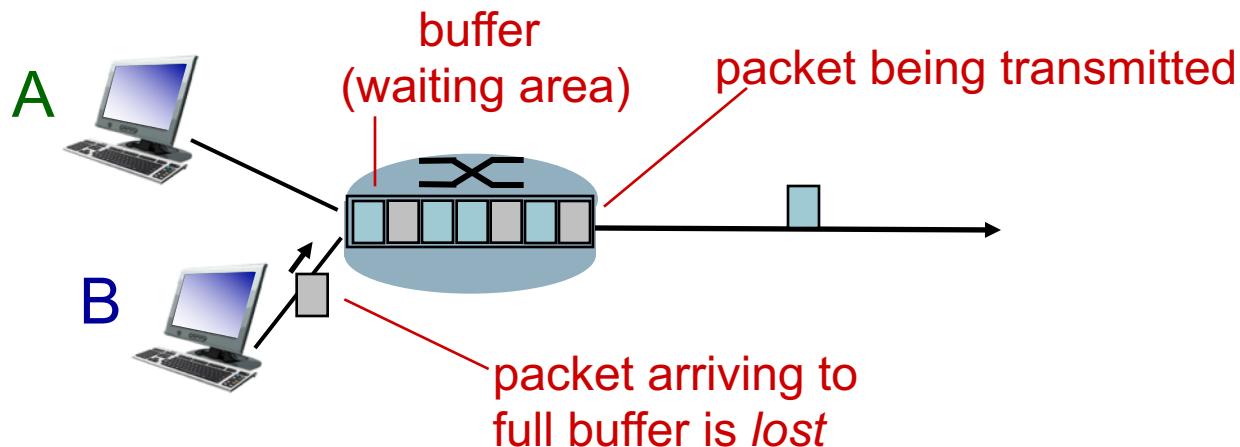
Queueing delay -- Congestion

- R : link bandwidth (bps)
 - L : packet length (bits)
 - a : average packet arrival rate
-
- ❖ $La/R \sim 0$: avg. queueing delay small
 - ❖ $La/R \rightarrow 1$: avg. queueing delay large
 - ❖ $La/R > 1$: more “work” arriving than can be serviced, average delay infinite!



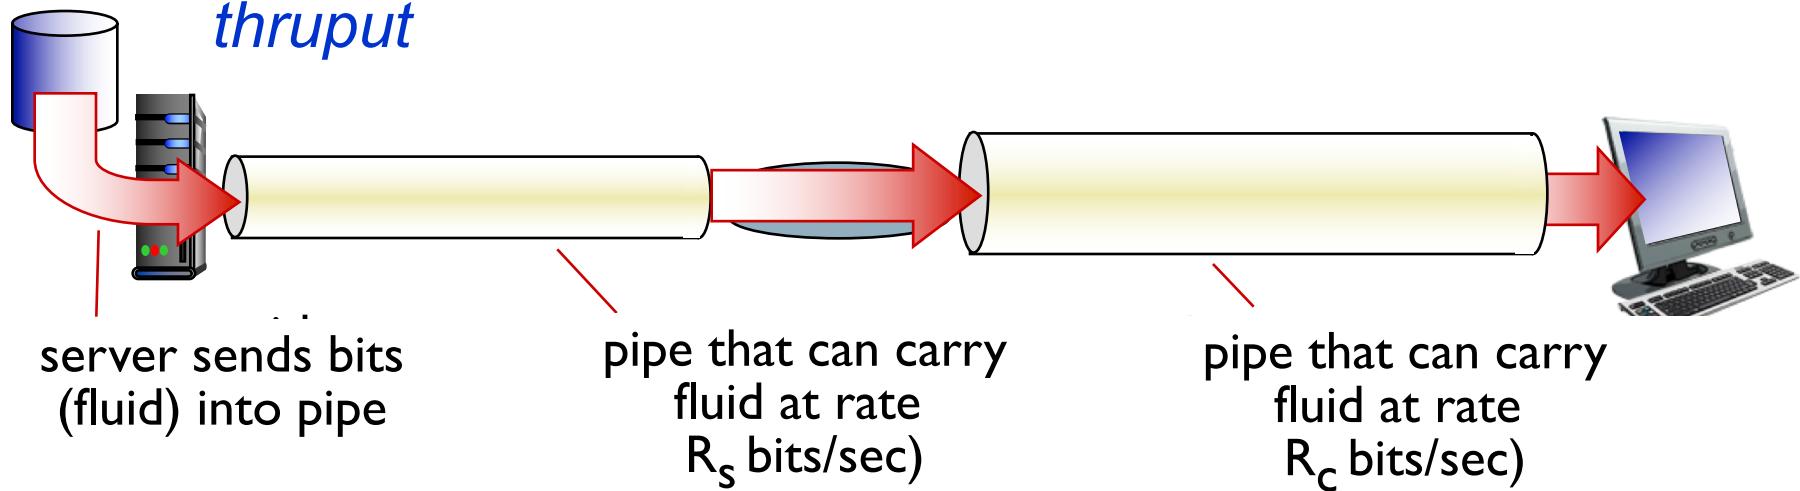
Packet loss

- queue (aka buffer) preceding link in buffer has finite capacity
- packet arriving to full queue dropped (aka lost)
- lost packet may be retransmitted by previous node, by source end system, or not at all

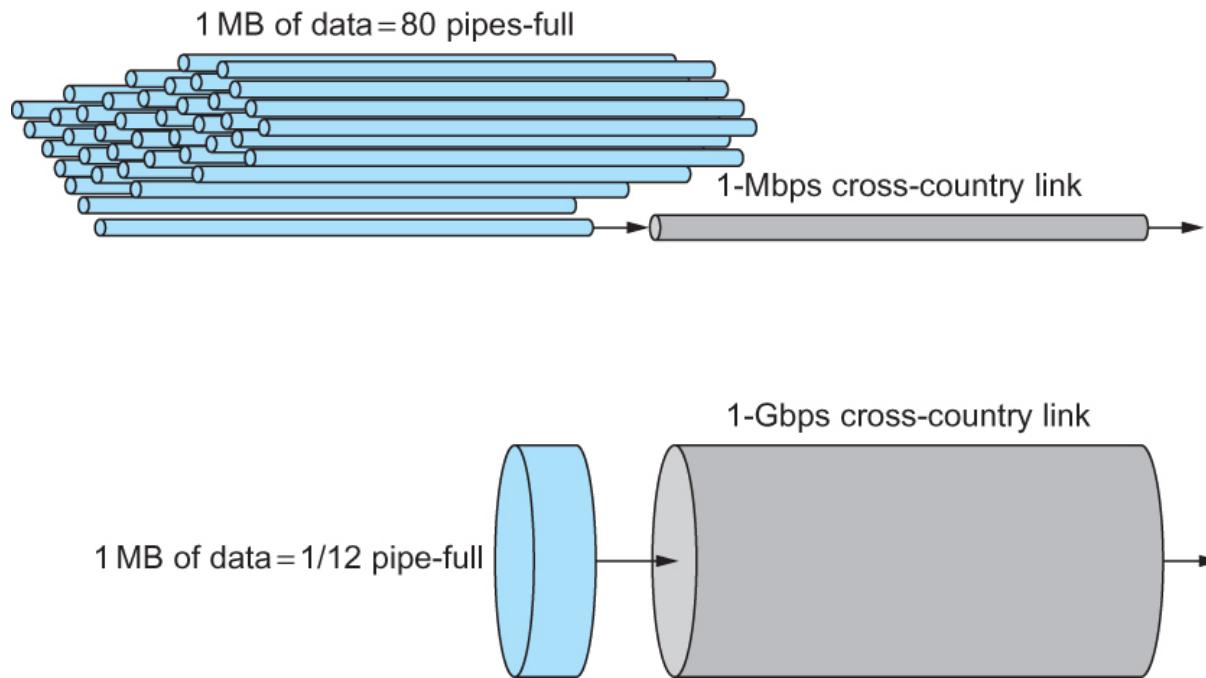


Throughput

- *throughput*: rate (bits/time unit) at which bits transferred between sender/receiver
 - *instantaneous*: rate at given point in time
 - *average*: rate over longer period of time
 - *bottleneck link on end to end path constrains throughput*



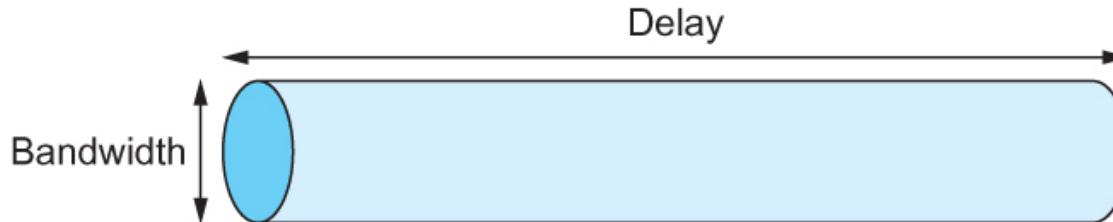
Relationship between bandwidth and latency



A 1-MB file would fill the 1-Mbps link 80 times,
but only fill the 1-Gbps link 1/12 of one time

Delay X Bandwidth

- We think the channel between a pair of processes as a hollow pipe
- Latency (delay) length of the pipe and bandwidth the width of the pipe
- Delay of 50 ms and bandwidth of 45 Mbps
 - ⇒ 50×10^{-3} seconds $\times 45 \times 10^6$ bits/second
 - ⇒ 2.25×10^6 bits = 280 KB data.



Network as a pipe

Delay X Bandwidth

- Relative importance of bandwidth and latency depends on application
 - For large file transfer, bandwidth is critical
 - For small messages (HTTP, NFS, etc.), latency is critical
 - Variance in latency (jitter) can also affect some applications (e.g., audio/video conferencing)
- How many bits the sender must/could transmit before the first bit arrives at the receiver
 - Takes another one-way latency to receive a response from the receiver
 - If the sender does not fill the pipe—send a whole delay \times bandwidth product's worth of data before it stops to wait for a signal—the sender will not fully utilize the network
 - Control travels over same network as data – latency impairs feedback which impairs thruput