

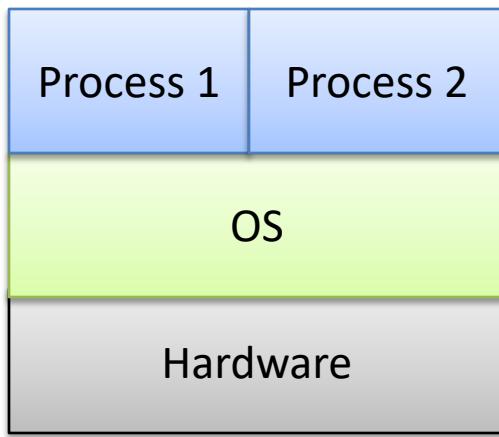
# **Virtualization security**

CS 5450

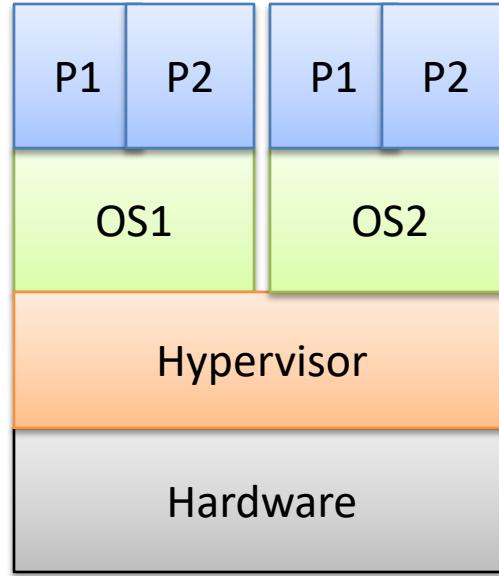
Tom Ristenpart

<https://rist.tech.cornell.edu>

# Virtualization



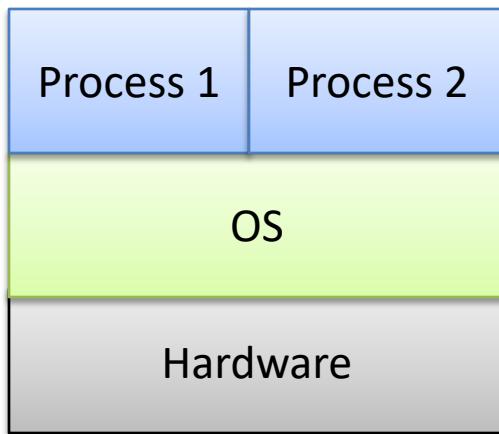
No virtualization



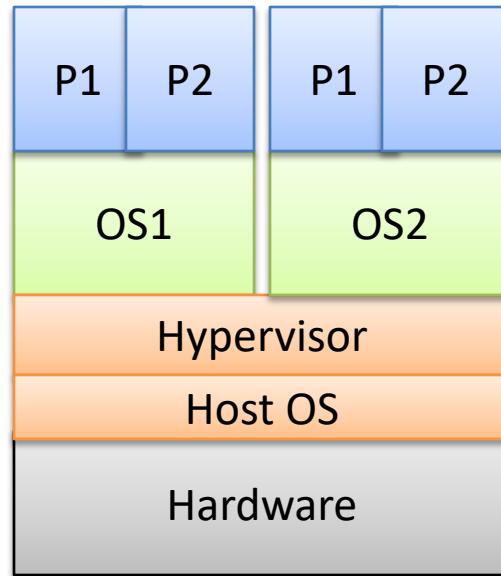
Full virtualization

Type-1: Hypervisor runs directly on hardware

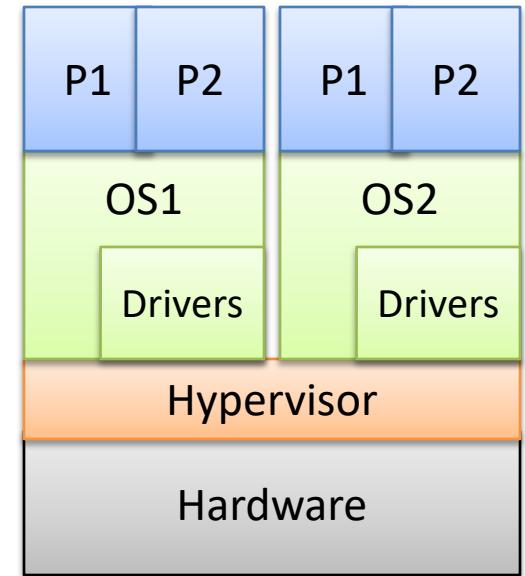
# Virtualization



No virtualization



Full virtualization



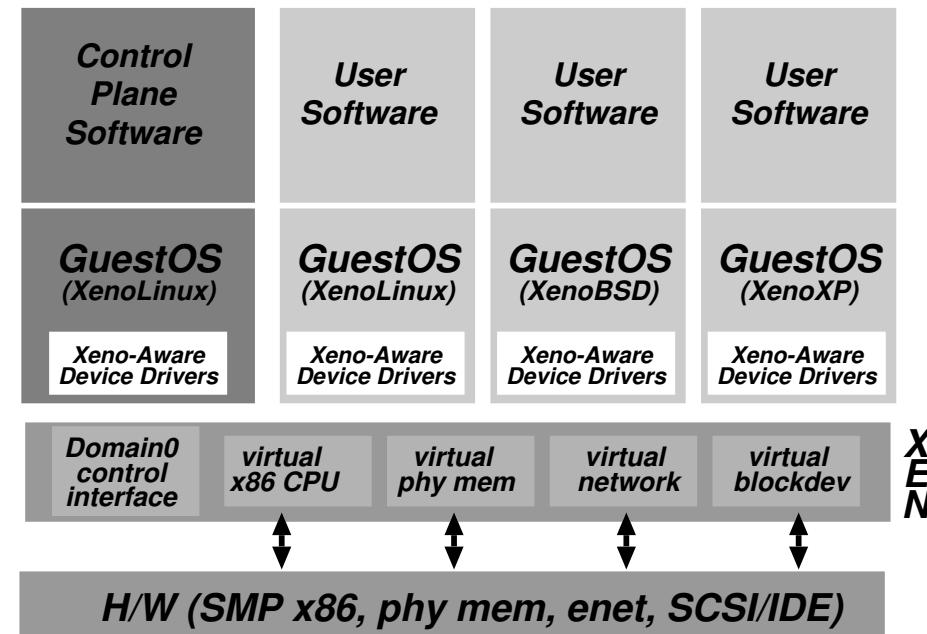
Paravirtualization

Type-1: Hypervisor runs directly on hardware  
Type-2: Hypervisor runs on host OS

# Xen



- 2003: academic paper
  - “Xen and the Art of Virtualization”
- Paravirtualization
  - Hypercalls vs system calls
  - Modified guest OS
  - Each guest given 1 or more VCPUs
- Why paravirtualization?



# Other modern VM solutions

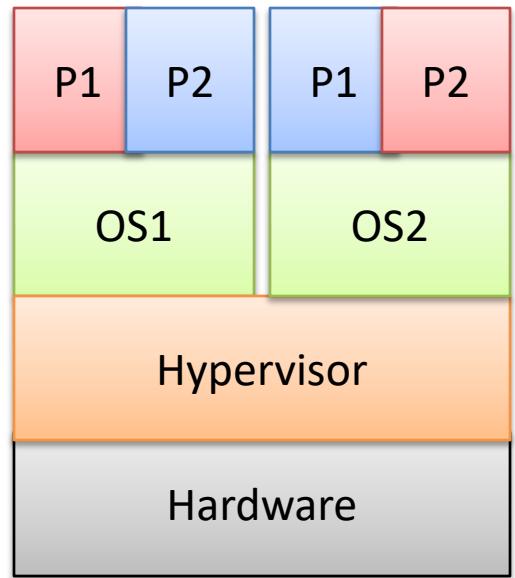
- VMWare
- Virtual Box
- KVM

# Example VM Use Cases

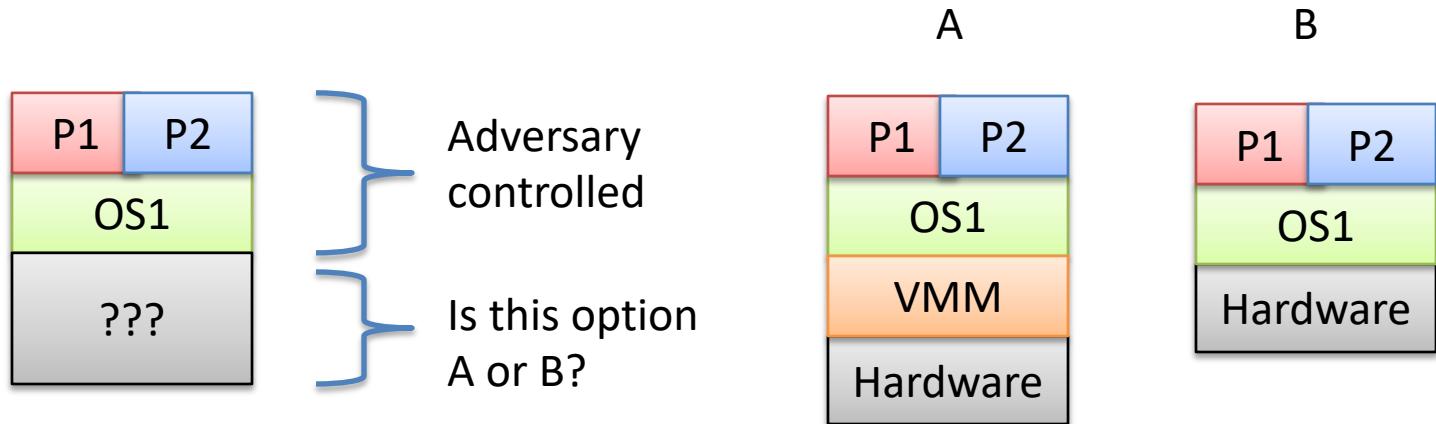
- Legacy support
- Development
- Server consolidation
- Cloud computing Infrastructure-as-a-Service
- Sandboxing / containment

# Study of malware

- Researchers use VMs to study malware
- Example of VM sandboxing
  - Hypervisor must contain malicious code
- Introspection (making sense of guest behavior from hypervisor) can be challenging
- How would you evade analysis as a malware writer?
  - split personalities



# VMM Transparency



- Adversary can detect if:
  - Paravirtualization
  - Logical discrepancies
    - Expected CPU behavior vs virtualized
    - Red pill (Store Interrupt Descriptor Table instr)
  - Timing discrepancies
    - Slower use of some resources

Garfinkel et al.  
“Compatibility  
is not transparency:  
VMM Detection  
Myths and Reality”

# Detection of VMWare

```
MOV EAX,564D5868 <-- "VMXh"
MOV EBX,0
MOV ECX,0A
MOV EDX,5658 <-- "VX"
IN EAX,DX <-- Check for VMWare
CMP EBX,564D5868
```

IN instruction used by VMWare  
to facilitate host-to-guest  
communication

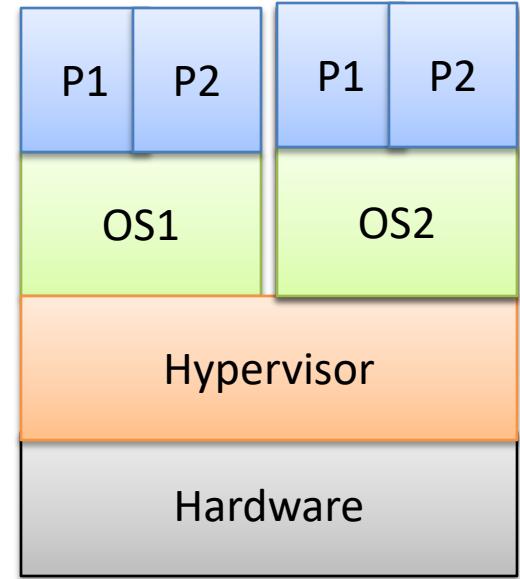
VMWare:  
    places VMXh in EBX  
Physical:  
    processor exception

From

[http://handlers.sans.org/tliston/ThwartingVMDetection\\_Liston\\_Skoudis.pdf](http://handlers.sans.org/tliston/ThwartingVMDetection_Liston_Skoudis.pdf)

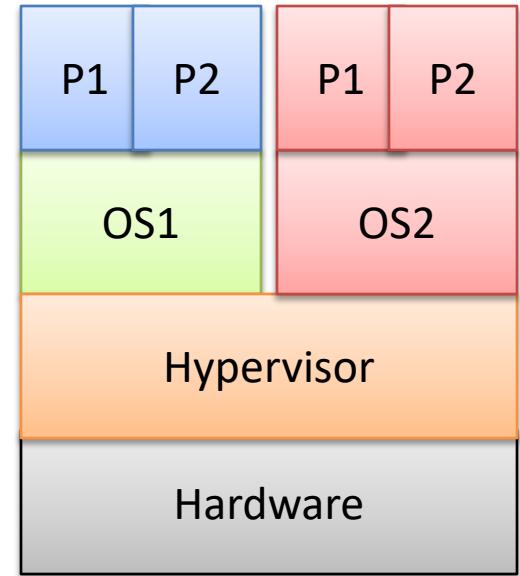
# Server consolidation

- Consolidation
  - Use VMs to optimize use of hardware
  - Pack as many VMs onto each server as possible
  - Turn off other servers
- Threat model?
  - Containment
  - Isolation
  - Assume guests are/can be compromised



# Violating containment

- Escape-from-VM
  - Vulnerability in VMM or host OS (e.g., Dom0)
  - Somewhat rare, but exist and very high value



## VMware vulnerability allows users to escape virtual environment

◦ By [Joab Jackson](#)   ◦ Feb 28, 2008

A new vulnerability found in some VMware products allows users to escape their virtual environments and muck about in the host operating system, penetration testing software firm Core Security Technologies [announced](#) earlier this week.

This vulnerability (CVE Name: CVE-2008-0923) could pose significant risks to enterprise users who are deploying VMware software as a secured environment.

NEWS

## Xen hypervisor faces third highly critical VM escape bug in 10 months

The Xen paravirtualization mode is proving to be a constant source of serious vulnerabilities, allowing attackers to escape from virtual machines

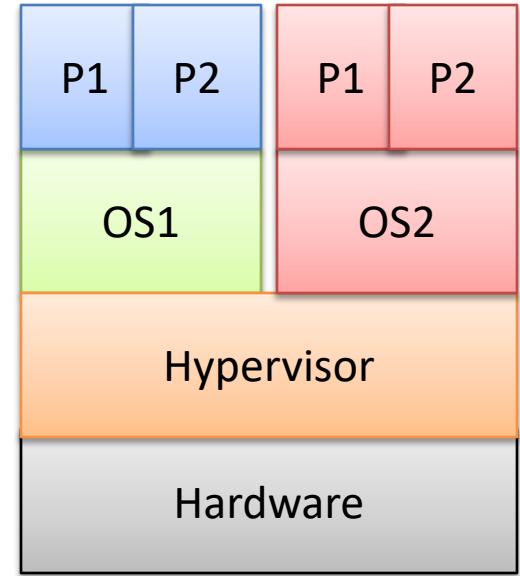
<https://blog.quarkslab.com/xen-exploitation-part-2-xsa-148-from-guest-to-host.html#id17>

# XSA-148 vulnerability

- Hypervisor must ensure guests can't read/write arbitrary physical memory addresses
- Xen uses ***direct paging***: guest VMs use virtual machine to physical address page tables
  - Must use hypercalls to update page tables
  - ***Security critical invariants must be enforced by hypervisor***
- XSA-148: Logic error in page table mapping checks, allowing guest to build page table mapping virtual addresses to ***any*** physical address
- ***Arbitrary read/write access to all of system memory***

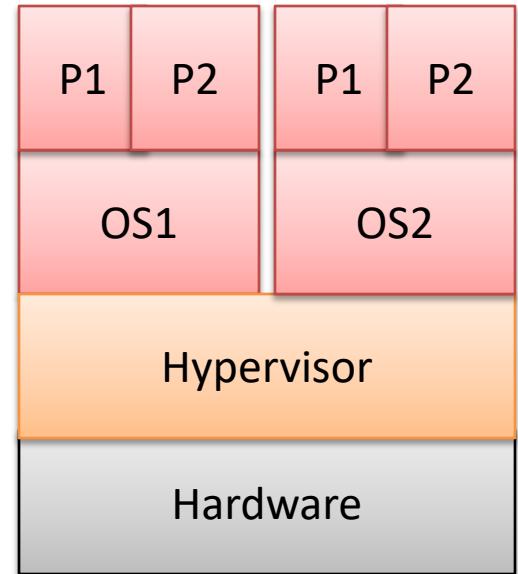
# Server consolidation

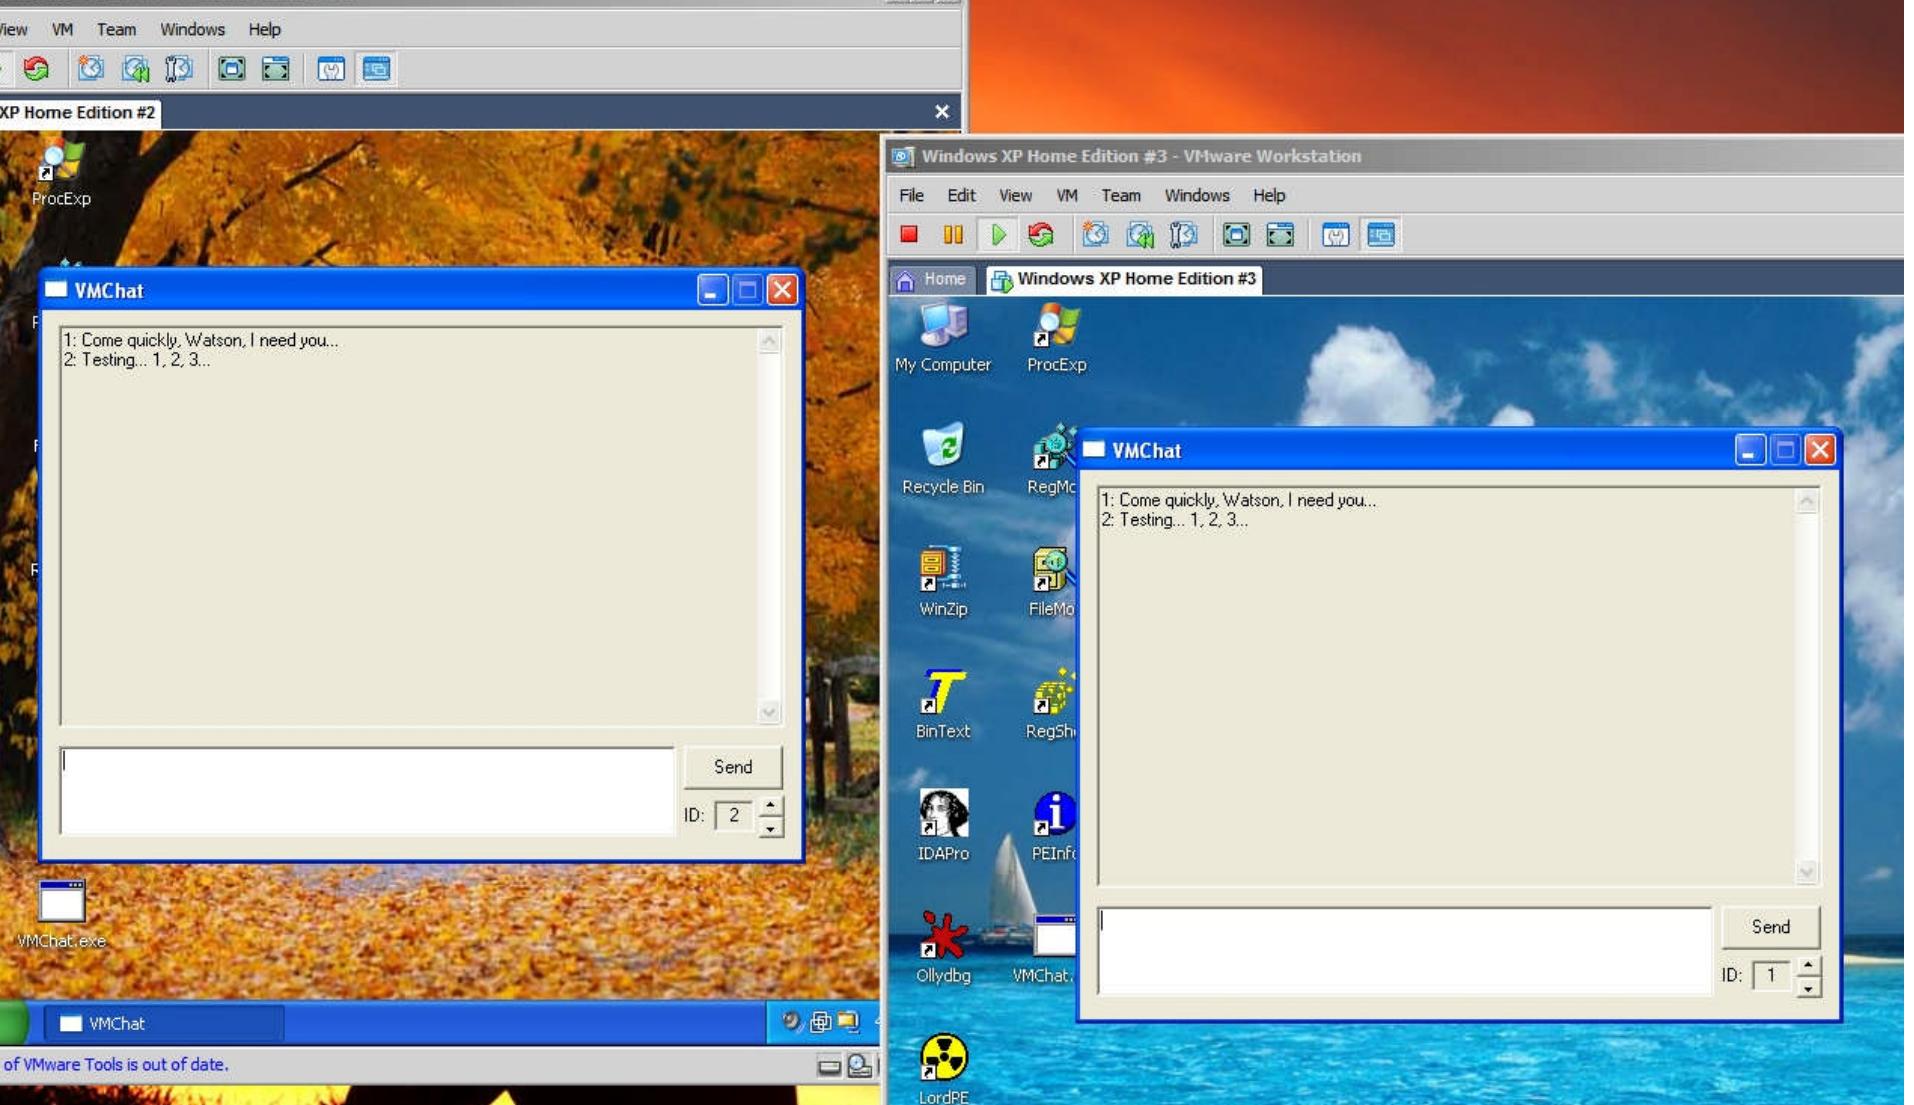
- Consolidation
  - Use VMs to optimize use of hardware
  - Pack as many VMs onto each server as possible
  - Turn off other servers
- Threat model?
  - Containment
  - Isolation ←
  - Assume guests are/can be compromised



# Violating isolation

- Covert channels between VMs circumvent access controls
  - Bugs in VMM
  - Side-effects of resource usage

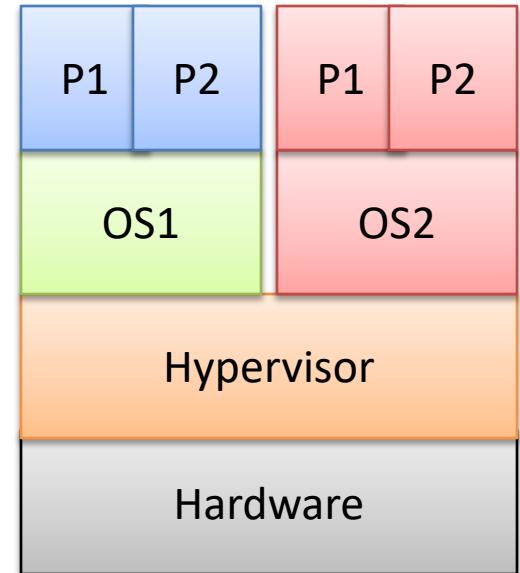




[http://handlers.sans.org/tliston/ThwartingVMDetection\\_Liston\\_Skoudis.pdf](http://handlers.sans.org/tliston/ThwartingVMDetection_Liston_Skoudis.pdf)

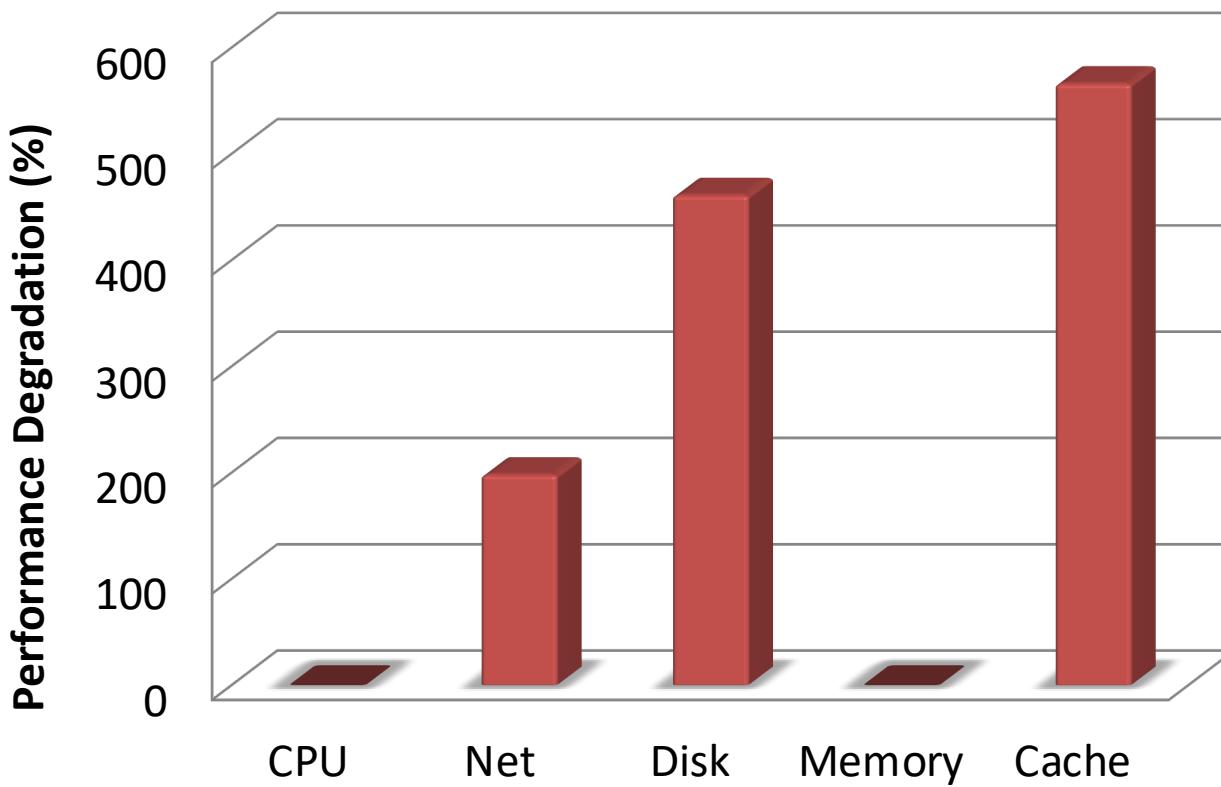
# Violating isolation

- Covert channels between VMs circumvent access controls
  - Bugs in VMM
  - Side-effects of resource usage
- Degradation-of-Service attacks
  - Guests might maliciously contend for resources
  - Xen scheduler vulnerability (<https://arxiv.org/abs/1103.0759>)



# Measuring Resource Contention

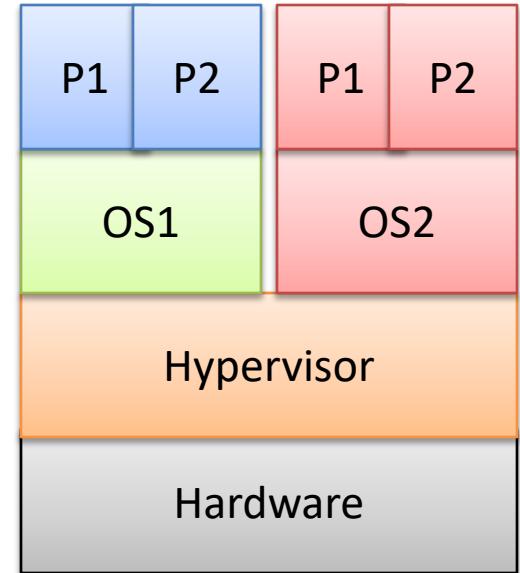
- Contention for the same resource



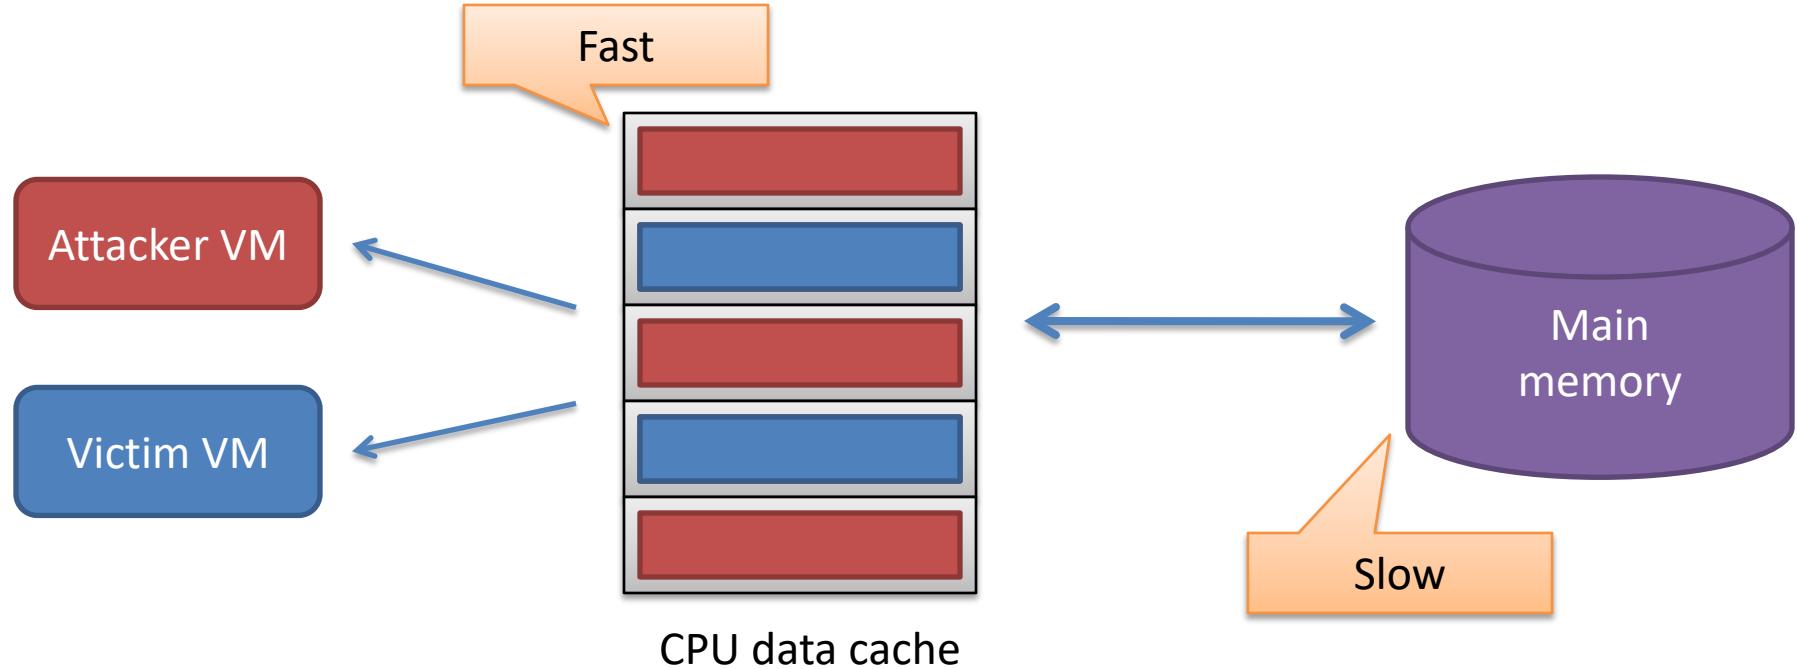
Local Xen Testbed	
Machine	Intel Xeon E5430, 2.66 Ghz
Packages	2, 2 cores per package
LLC Size	6MB per package

# Violating isolation

- Covert channels between VMs circumvent access controls
  - Bugs in VMM
  - Side-effects of resource usage
- Degradation-of-Service attacks
  - Guests might maliciously contend for resources
  - Xen scheduler vulnerability (<https://arxiv.org/abs/1103.0759>)
- Side channels
  - Spy on other guest via shared resources



# Cross-VM side channels using CPU cache contention

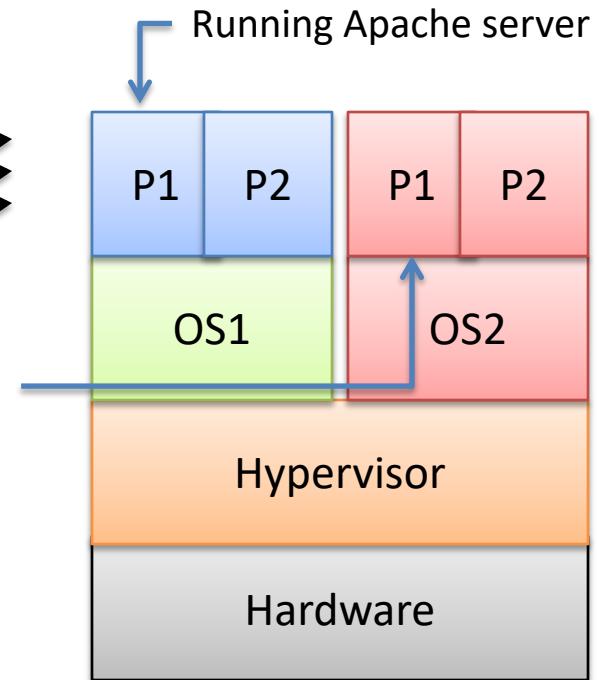
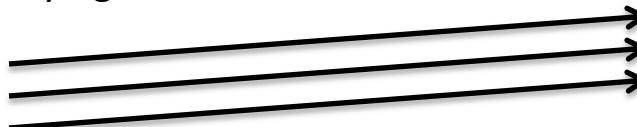


- 1) Read in a large array (fill CPU cache with attacker data)
- 2) Busy loop (allow victim to run)
- 3) Measure time to read large array (the load measurement)

# Cache load is correlated with Apache traffic volume

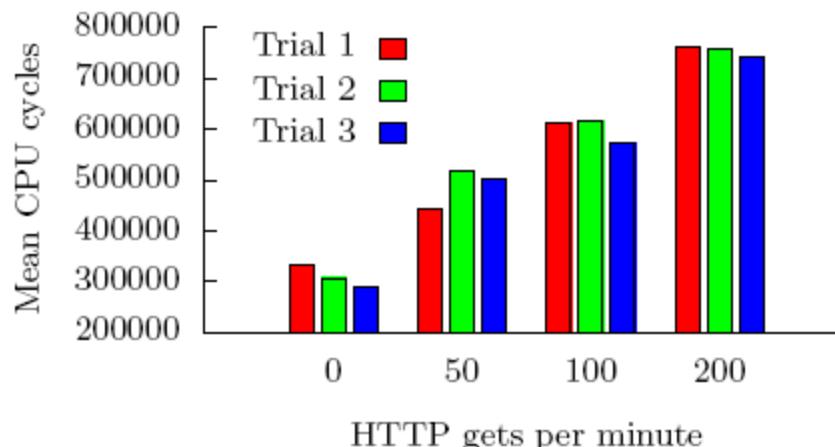


Varying rates of web traffic



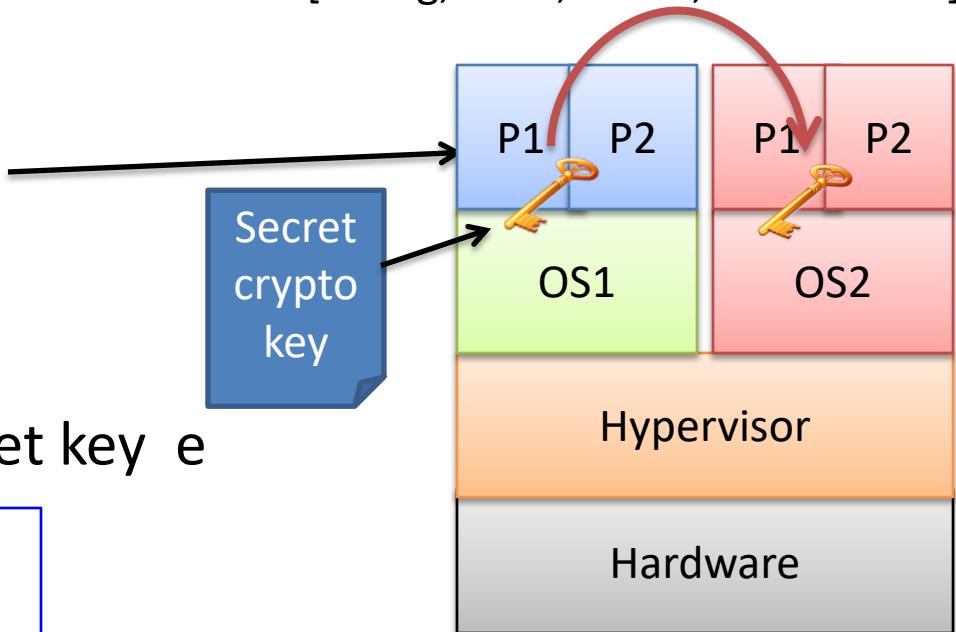
3 trials with 1 pair of EC2 instances:

1,000 cache load measurements during  
0, 50, 100, or 200 **HTTP gets** (3 Mbyte page) per  
minute for ~1.5 mins



# Cross-VM cryptographic side-channel attacks

[Zhang, Juels, Reiter, R. – CCS '12]



Target is 4096-bit ElGamal secret key  $e$

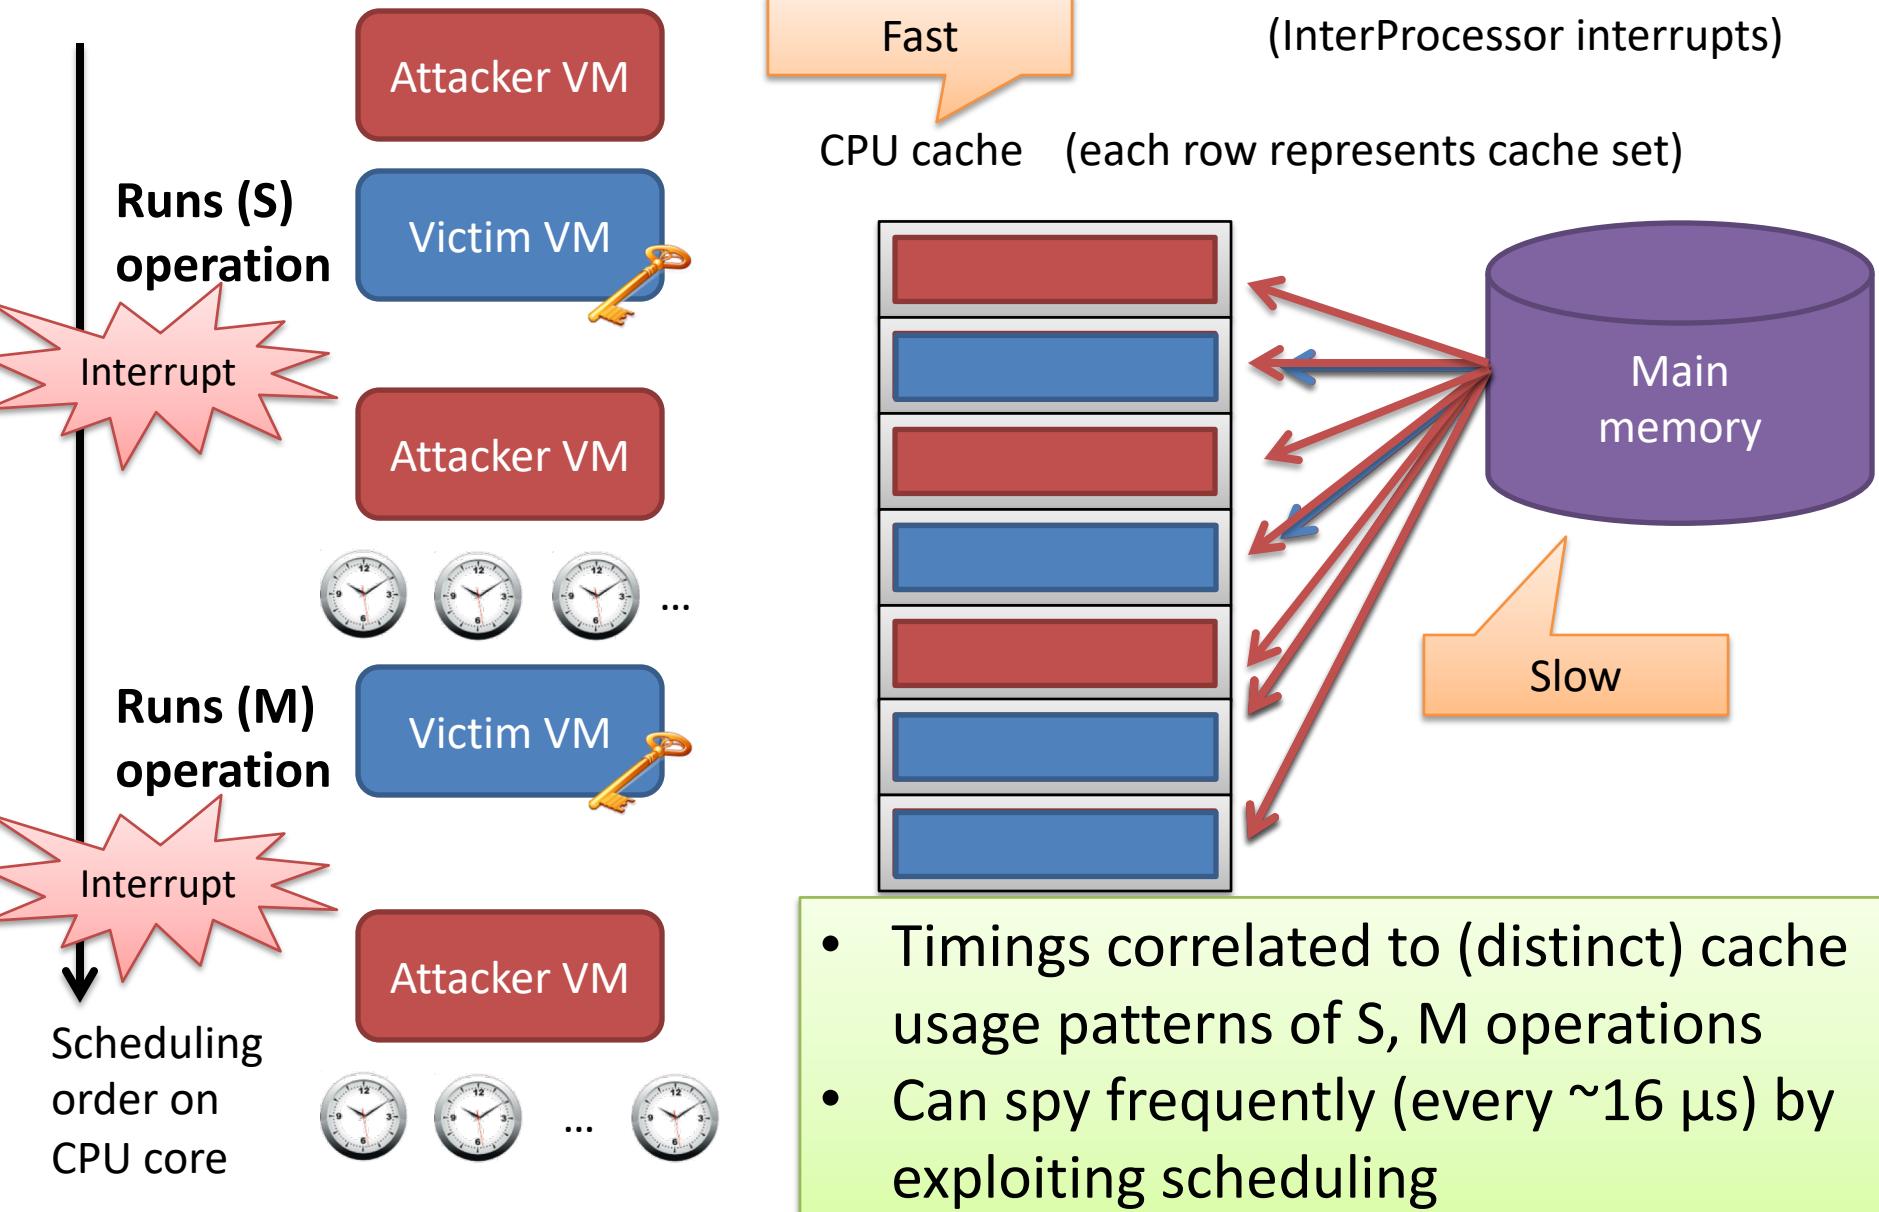
## Modular Exponentiation ( $x, e, N$ ):

```
let  $e_n \dots e_1$  be the bits of  $e$ 
 $y \leftarrow 1$ 
for  $e_i$  in  $\{e_n \dots e_1\}$ 
     $y \leftarrow \text{Square}(y)$           (S)
     $y \leftarrow \text{Reduce}(y, N)$       (R)
    if  $e_i = 1$  then
         $y \leftarrow \text{Multi}(y, x)$     (M)
         $y \leftarrow \text{Reduce}(y, N)$       (R)
return  $y$  //  $y = x^e \bmod N$ 
```

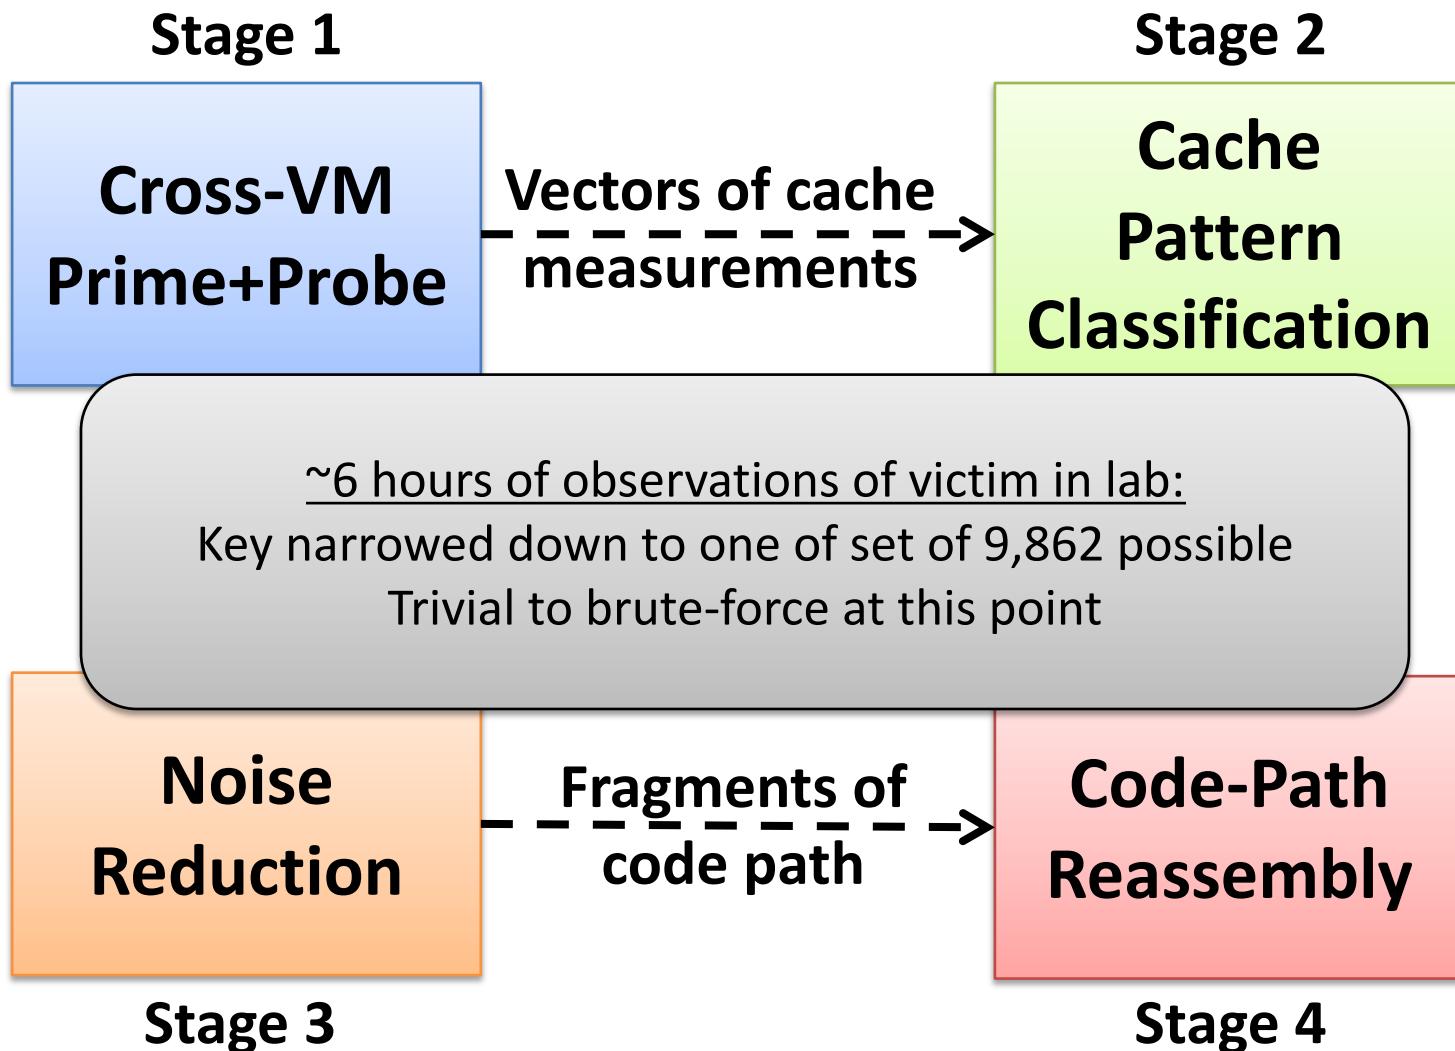
$e_i = 1 \rightarrow \text{"SRMR"}$   
 $e_i = 0 \rightarrow \text{"SR"}$

Sequence of function calls  
reveals secret key

# Prime+Probe protocol using IPIs



# Full attack pipeline

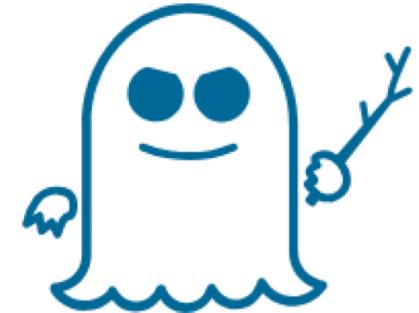


# Cross-VM side channel attacks

- Active research on new attacks and defenses for side-channels
- Most recent attacks can steal secrets quickly
  - Almost exclusively in lab settings
  - One paper demonstrated RSA secret key recovery across VMs on EC2 [Inci et al. 2016]
- Unclear if used by real attackers in the wild

# New class of vulnerabilities

- Meltdown/Spectre CPU vulnerabilities discovered 2017, made public early 2018
- Memory access controls bypassed due to how speculative execution works
- Allow reading victim memory



# Speculative execution

Modern CPUs allow out-of-order execution of instructions to speed things up

Use branch prediction to attempt to guess outcome of conditional statements

- (1)     if ( $x < \text{array1\_size}$ )
- (2)          $y = \text{array2}[\text{array1}[x] * 4096];$

## Normal execution:

Fetch `array1_size` from memory, compare to  $x$   
If condition holds, execute (2), otherwise do not

# Speculative execution

Modern CPUs allow out-of-order execution of instructions to speed things up

Use branch prediction to attempt to guess outcome of conditional statements

```
(1)    if (x < array1_size)  
(2)        y = array2[array1[x] * 4096];
```

## Speculative execution:

- Say `array1_size` must be fetched from memory (very slow!)
- Checkpoint state of CPU registers
- If branch prediction guesses that  $x < \text{array1\_size}$  then run line (2) before determining  $x < \text{array1\_size}$
- Now check if  $x < \text{array1\_size}$ , commit if so, otherwise rollback state



Line (2) gets executed as ***transient instructions*** whether or not condition is true

# New class of vulnerabilities

Say this snippet of code in hypercall path

x is controlled by malicious attacker

```
(1)    if (x < array1_size)  
(2)        y = array2[array1[x] * 4096];
```

array1[x] secret value  
(e.g., one byte of a  
cryptographic secret)

Transient execution fetches

k = array1[x]  
and array2[k\*4096]

Here k is a secret!

Attacker checks what value  
of k has

array2[k\*4096]  
loaded in cache via side-  
channel

# New class of vulnerabilities

Say this snippet of code in hypercall path

x is controlled by malicious attacker

```
(1)    if (x < array1_size)  
(2)        y = array2[array1[x] * 4096];
```

array1[x] secret value  
(e.g., one byte of a  
cryptographic secret)

Transient execution fetches

$k = \text{array1}[x]$

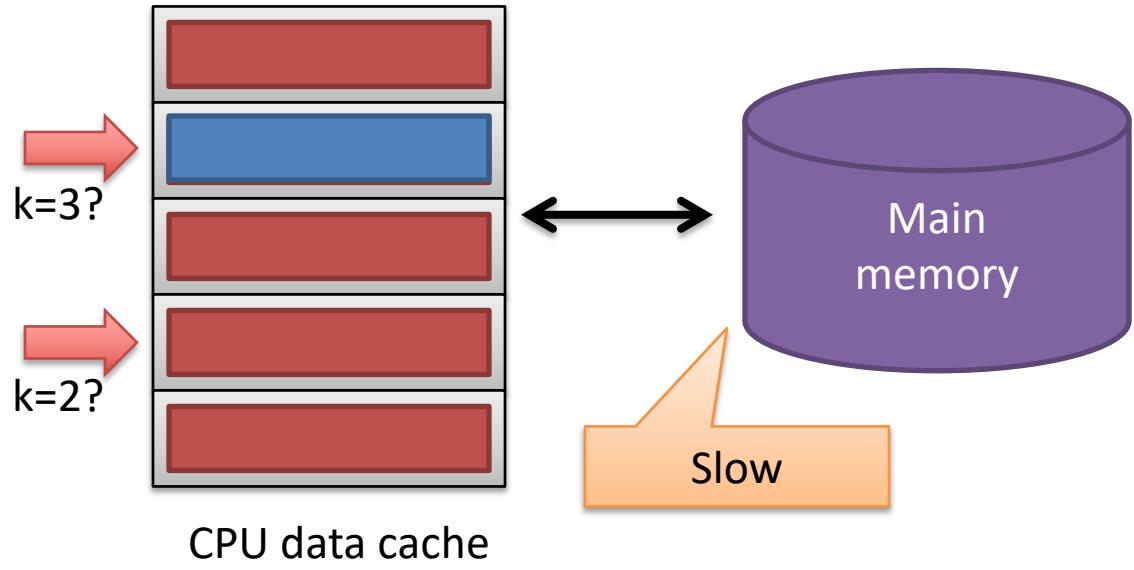
and  $\text{array2}[k*4096]$

Here k is a secret!

Attacker checks what value  
of k has

$\text{array2}[k*4096]$

loaded in cache via side-  
channel



# Fallout of Meltdown/Spectre

- Software mitigations for some exploits
- VM vendors released patches
- Cloud providers had to roll out patches quickly
- Intel had to ship new microcode for CPUs
  - Performance degradation serious concern
- 3 class-action lawsuits against Intel

# VM Containment & Isolation Lessons

- VMMs provide better containment and isolation than:
  - Processes, containers (e.g., Docker)
- But still not perfect:
  - VMM transparency
  - Vulnerabilities in hypervisor and CPU
  - Side channels exist
- Securing guest OS and host OS still very important for defense-in-depth

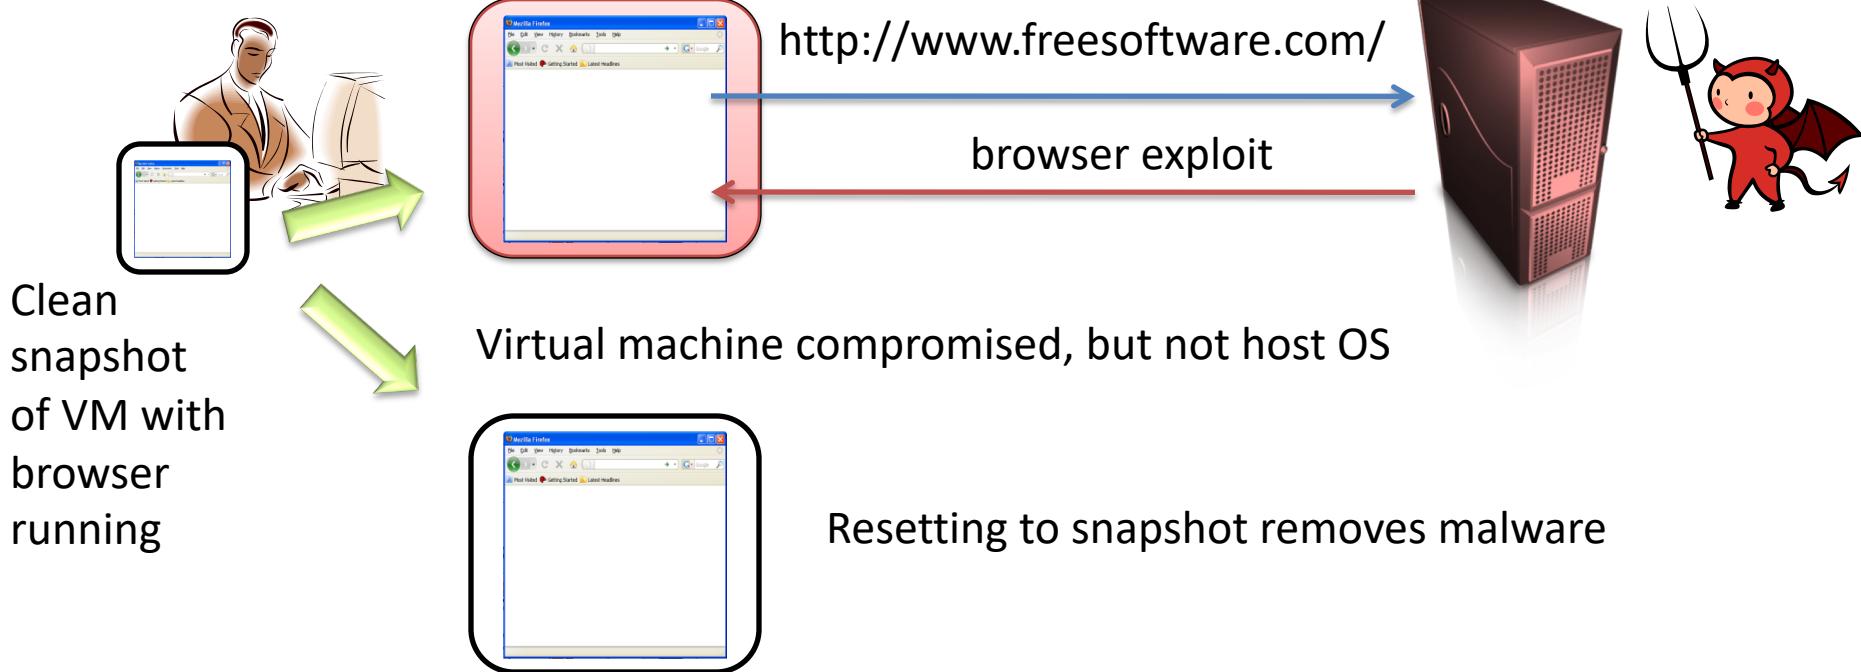
# Virtual Machine lifecycle management

- Snapshots
  - Volume snapshot / checkpoint
    - persistent storage of VM
    - must boot from storage when resuming snapshot
  - Full snapshot
    - persistent storage and ephemeral storage (memory, register states, caches, etc.)
    - start/resume in between (essentially) arbitrary instructions
- VM image is a file that stores a snapshot

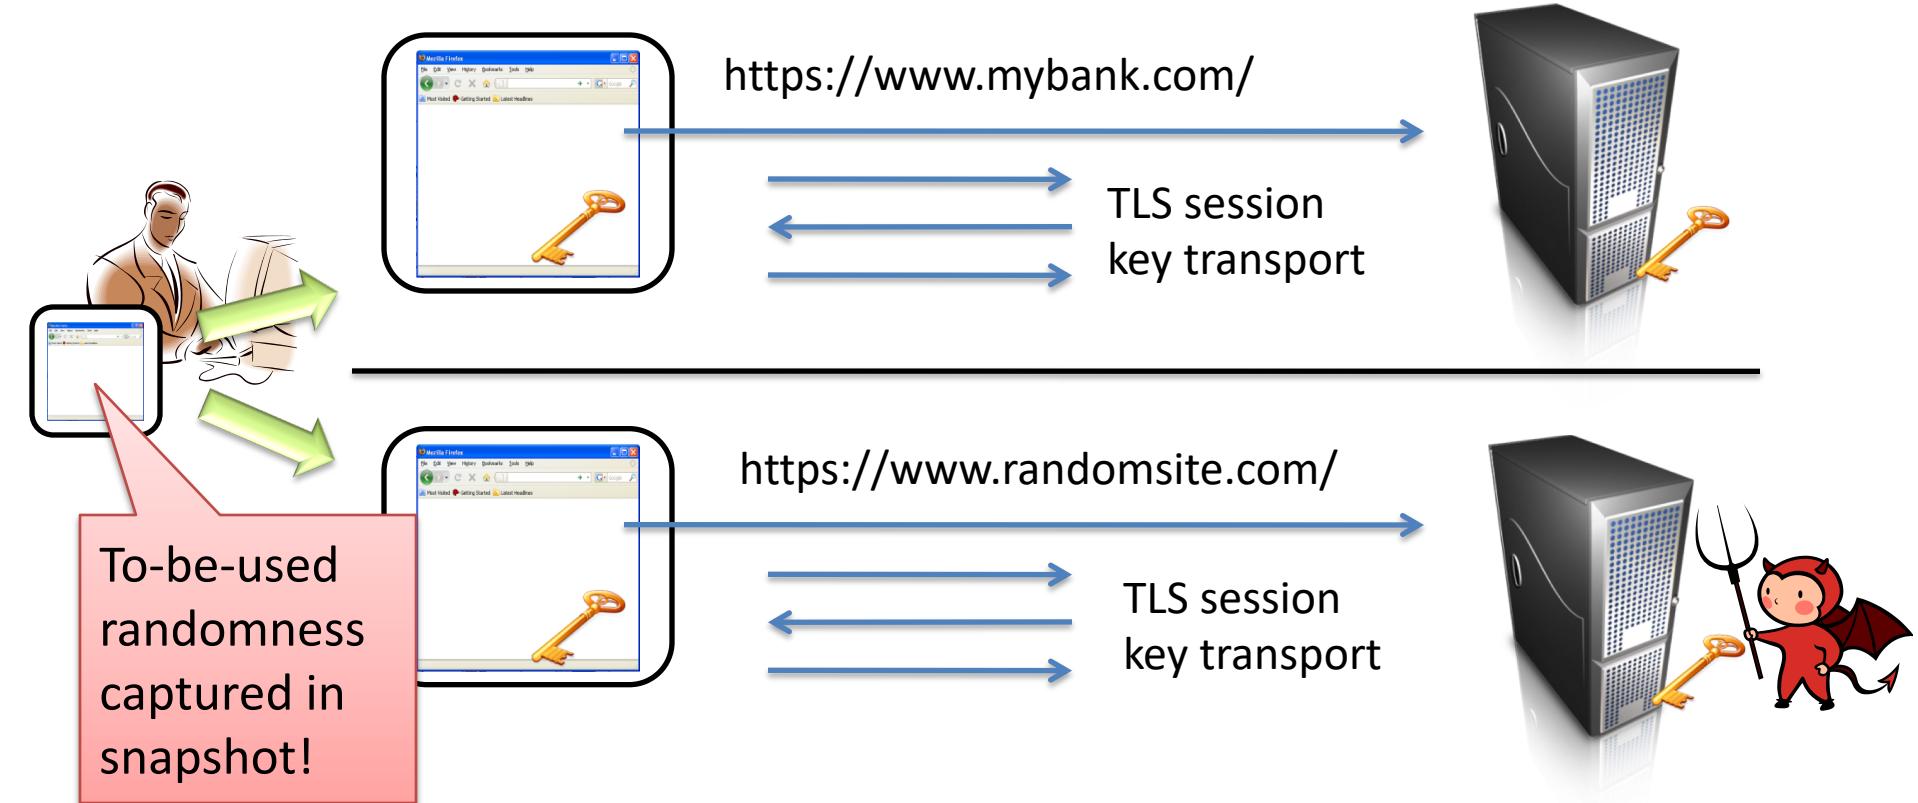
# Virtual machines and secure browsing

**"Protect Against Adware and Spyware:** Users protect their PCs against adware, spyware and other malware while browsing the Internet with Firefox in a virtual machine."

[<http://www.vmware.com/company/news/releases/player.html>]



# Virtual machine resets lead to RNG failures



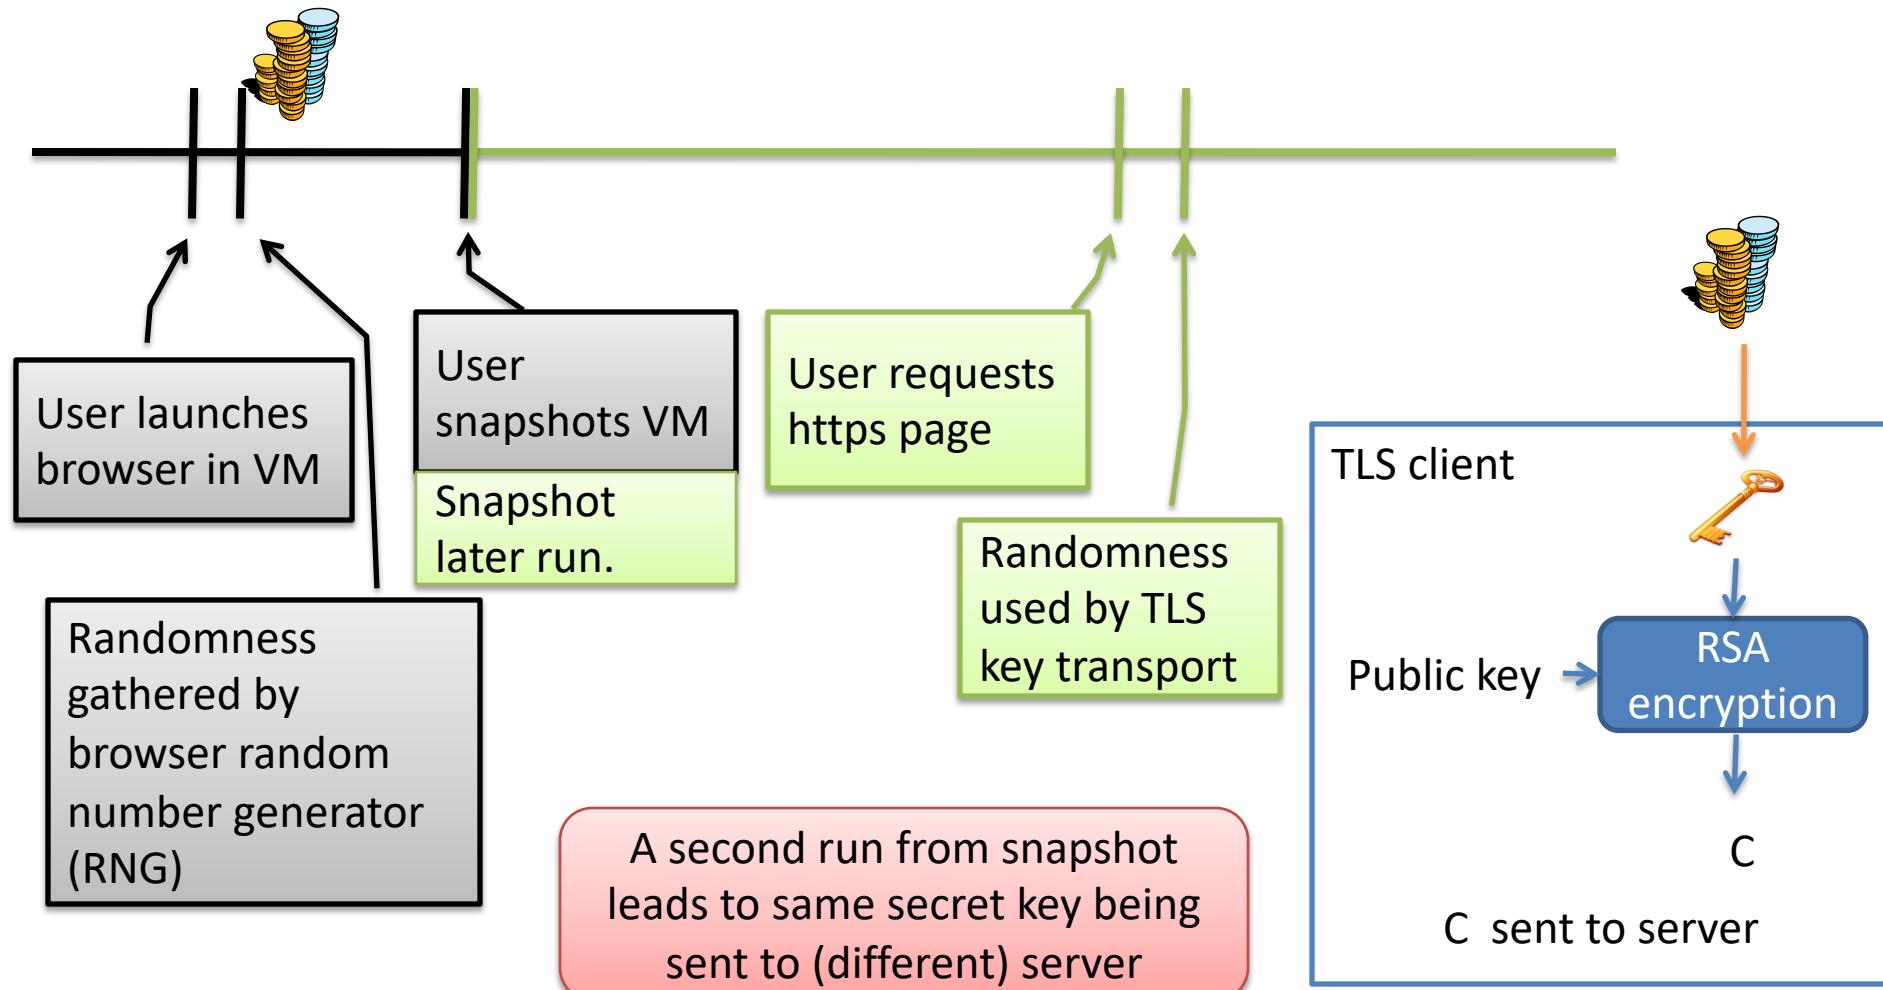
## 2010 vulnerabilities:

Some versions of **Firefox**, **Chrome** allow session compromise attacks

Apache mod\_ssl TLS server:  
server's secret key can be stolen!

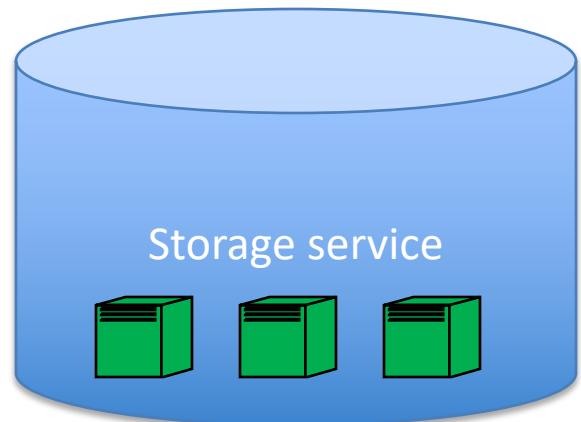


### A logical timeline of events

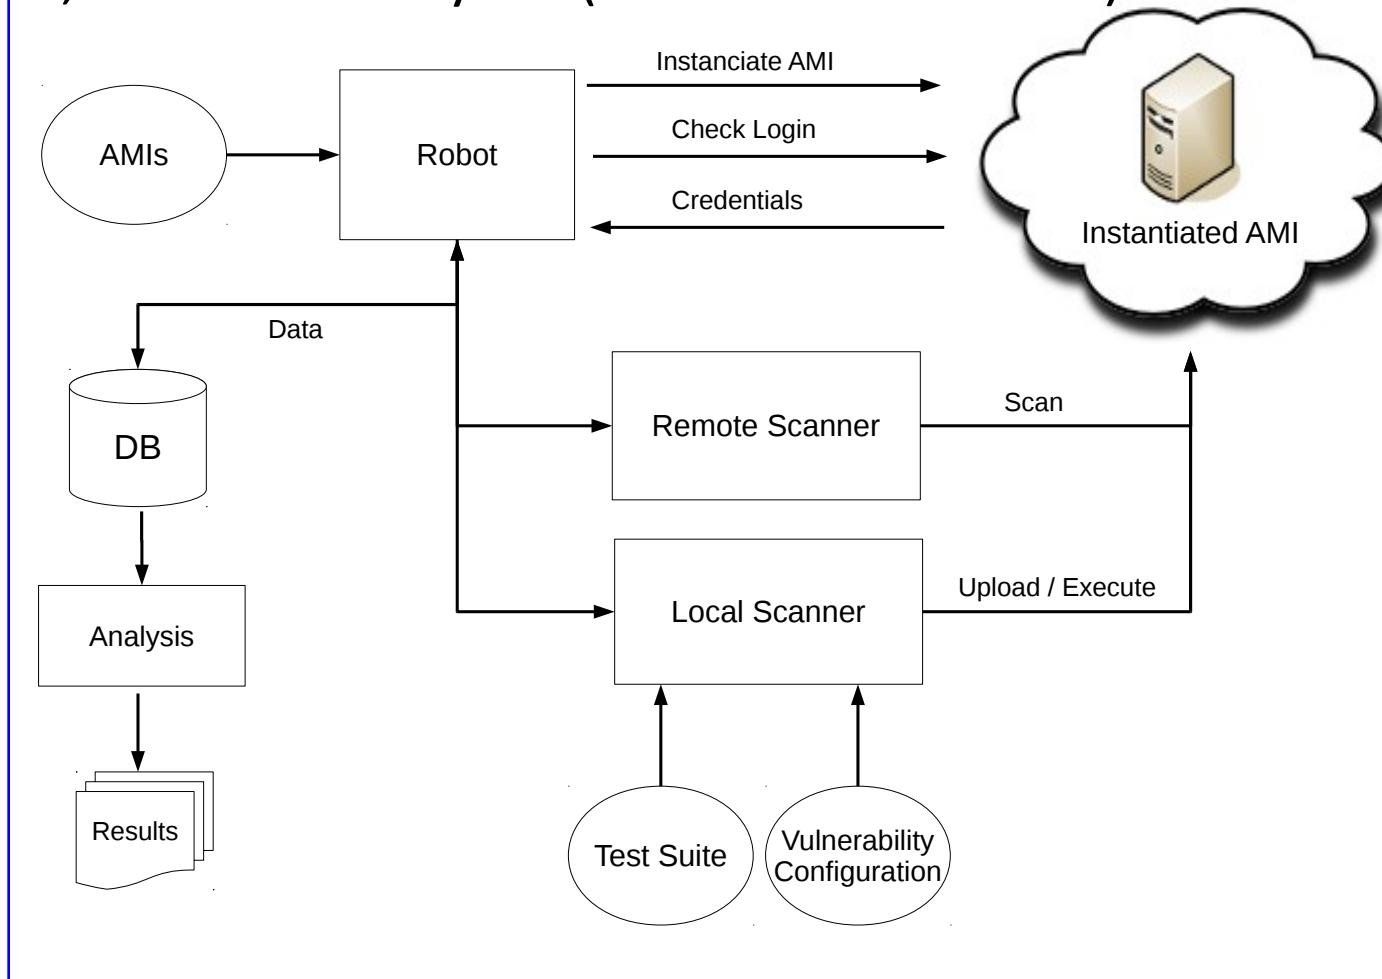


# Amazon Machine Images (AMIs)

- Users set up volume snapshots / checkpoints that can then be run on the Elastic Compute Cloud (EC2)
- Can be marked as public and anyone can use your AMI

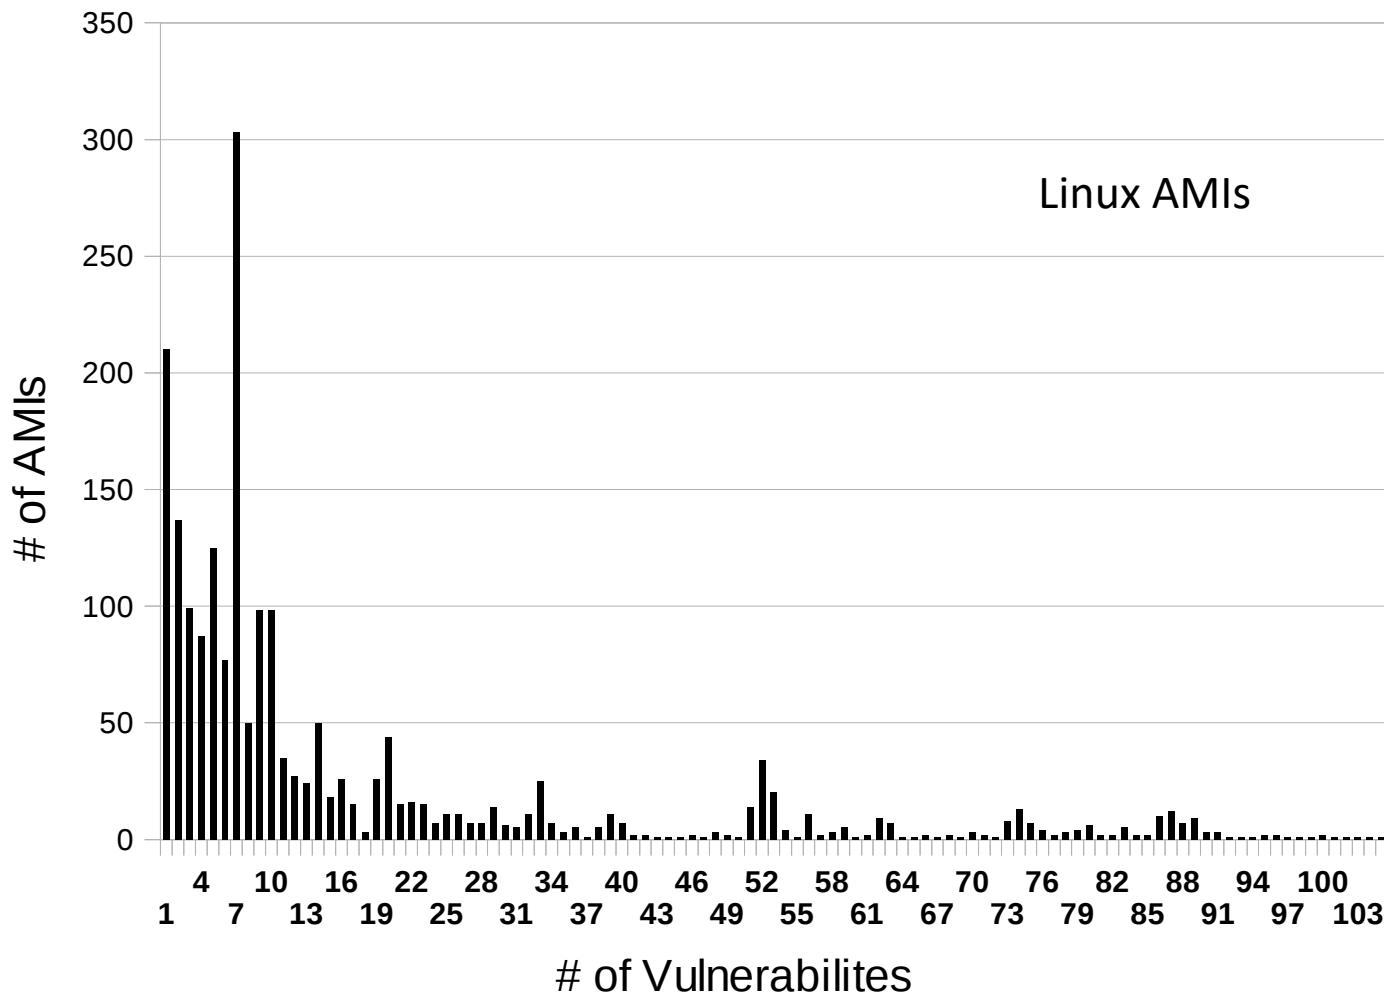


# 5,303 AMIs analyzed (Linux and Windows)



Baldazzi et al. "A Security Analysis of Amazon's Elastic Compute Cloud Service – Long Version –", 2011

See also Bugiel et al., "AmazonIA: When Elasticity Snaps Back", 2011



Also: Malware found on a couple AMIs

# Baldazzi et al. analysis

- Backdoors
  - AMIs include SSH public keys within `authorized_keys`
  - Password-based backdoors

	East	West	EU	Asia	Total
AMIs (%)	34.8	8.4	9.8	6.3	21.8
With Passwd	67	10	22	2	101
With SSH keys	794	53	86	32	965
With Both	71	6	9	4	90
Superuser Priv.	783	57	105	26	971
User Priv.	149	12	12	12	185

**Table 2: Left credentials per AMI**

# Baldazzi et al. analysis

- Credentials for other systems
  - AWS secret keys (to control EC2 services of an account): 67 found
  - Passwords / secret keys for other systems: 56 found

Finding	Total	Image	Remote
Amazon RDS	4	0	4
dDNS	1	0	1
SQL	7	6	1
MySQL	58	45	13
WebApp	3	2	1
VNC	1	1	0
Total	74	54	20

Table 3: Credentials in history files

# Baldazzi et al. analysis

- Deleted files
  - One AMI creation method does block-level copying

Type	#
Home files (/home, /root)	33,011
Images (min. 800x600)	1,085
Microsoft Office documents	336
Amazon AWS certificates and access keys	293
SSH private keys	232
PGP/GPG private keys	151
PDF documents	141
Password file (/etc/shadow)	106

**Table 5: Recovered data from deleted files**

# Response

<http://www.forbes.com/sites/andygreenberg/2011/11/08/>

researchers-find-amazon-cloud-servers-teeming-with-backdoors-and-other-peoples-data/

- Amazon notified customers with vulnerable AMIs
- Made private AMIs of non-responsive customers
- New tutorials for bundling systems

# Similar problem: git AWS exposures

- Common bad practice to include AWS credentials in source code
  - Gets checked into public git repos
- Criminals and Amazon now regularly scan github and other sites for exposed credentials
- <https://github.com/aws-labs/git-secrets>

# VM management lessons

- New software management practices & tools needed for new software deployment models
- Use tools to help prevent credential exposure, use provider's credential management tools (AWS has many)

